

LLVM* IR TUTORIAL - PHIS, GEPS AND OTHER THINGS, OH MY!

LLVM Developers Conference – Brussels 2019

Vince Bridgers

Felipe de Azevedo Piovezan

About this tutorial

Assume **no** previous Intermediate Representation (IR) knowledge.

But this is not a lecture about compiler theory!

After the tutorial, you should:

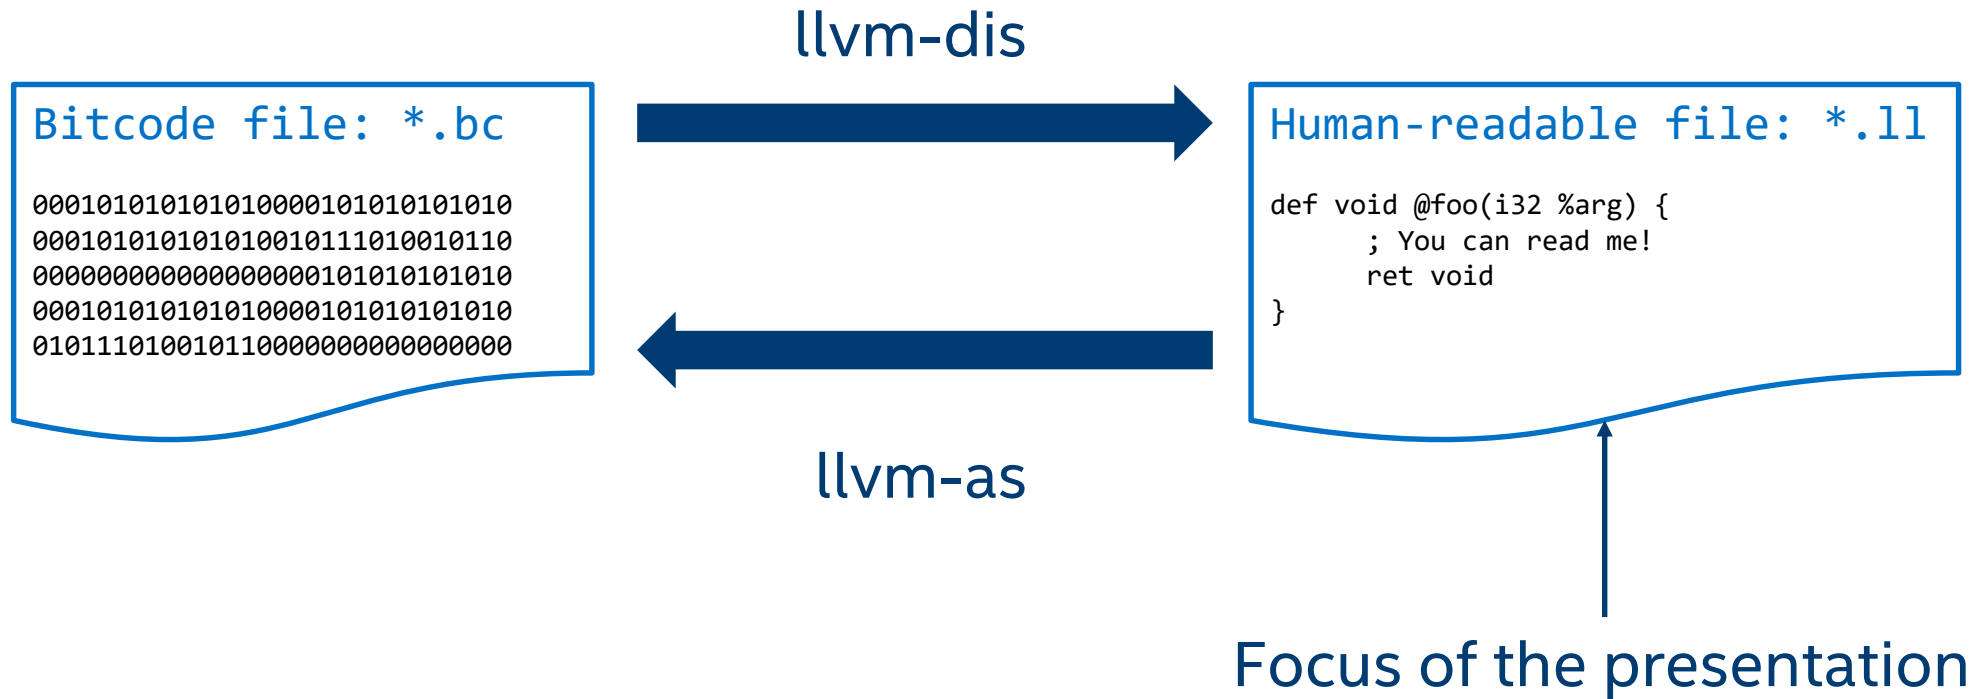
- Understand common LLVM tools.
- Be able to write simple IR.
- Be able to understand the [language reference](#).
 - Use it to inspect compiler-generated IR.

What is the LLVM IR?

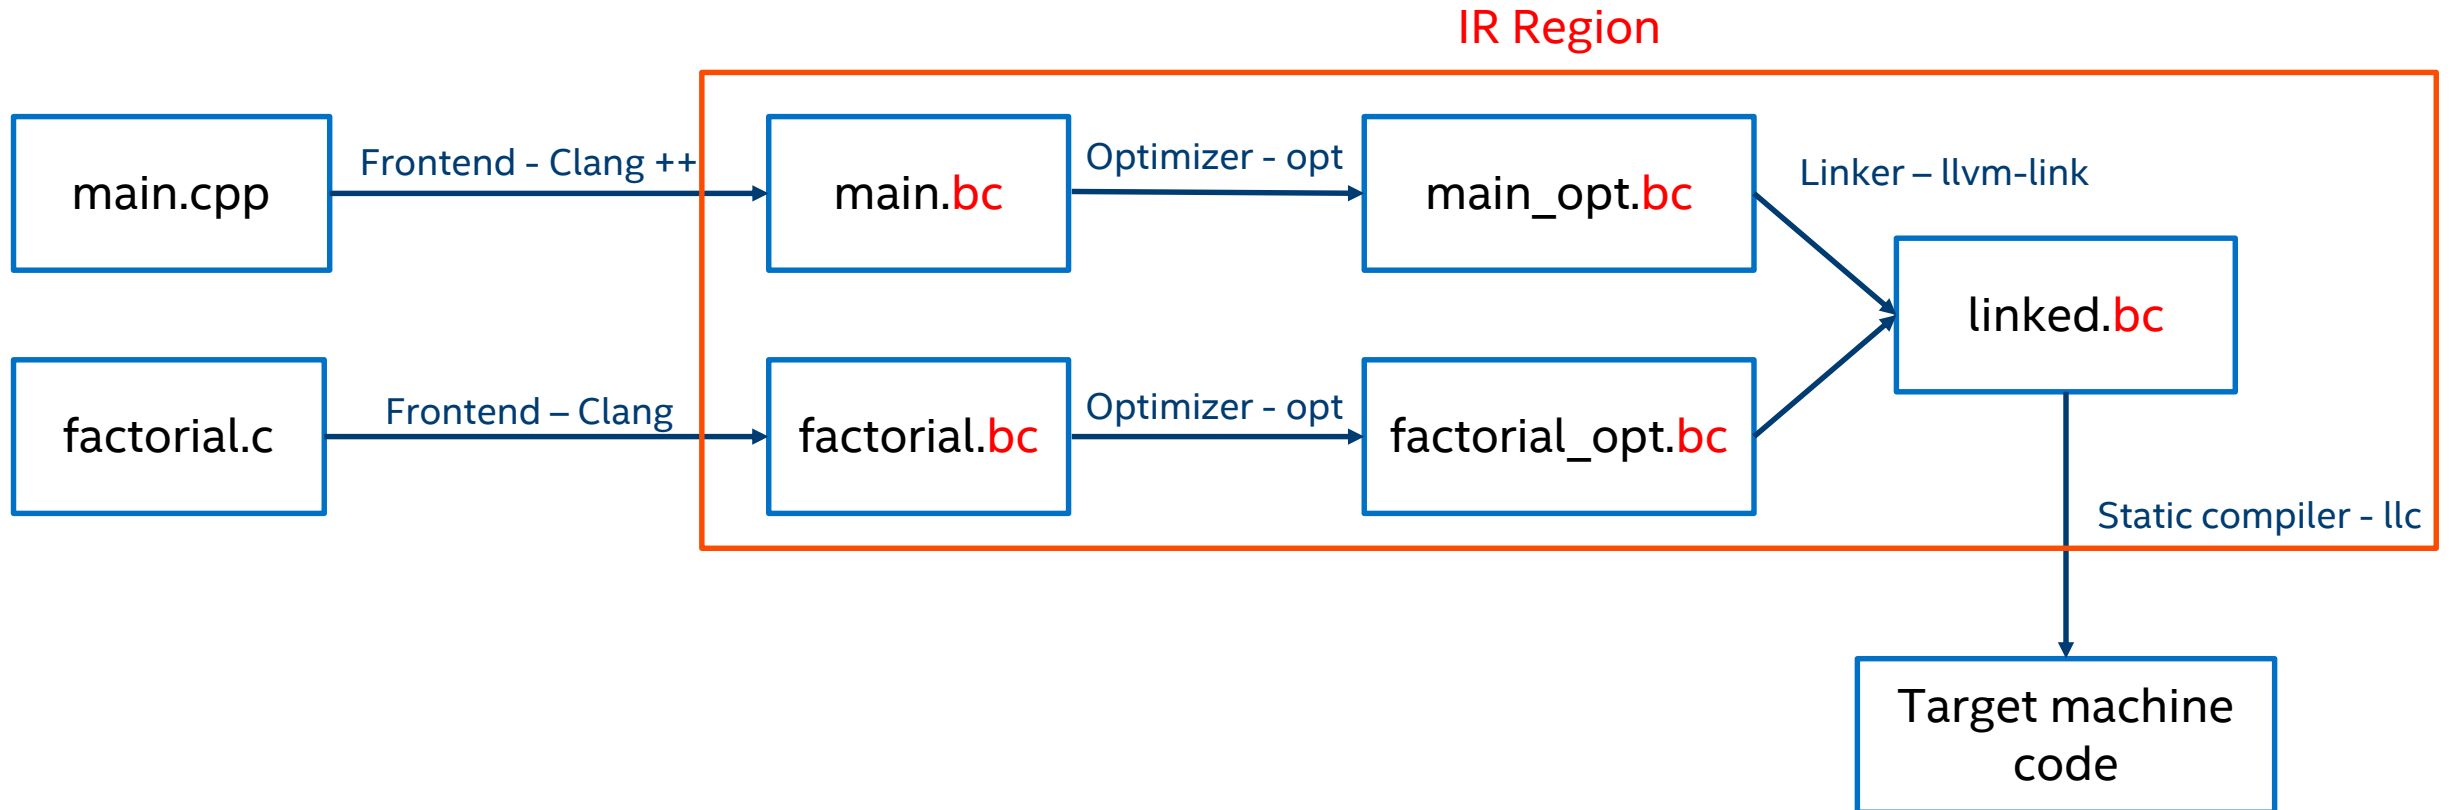
The LLVM Intermediate Representation:

- Is a low level programming language
 - RISC-like instruction set
- ... while being able to represent high-level ideas.
 - i.e. high-level languages can map cleanly to IR.
- Enables efficient code optimization

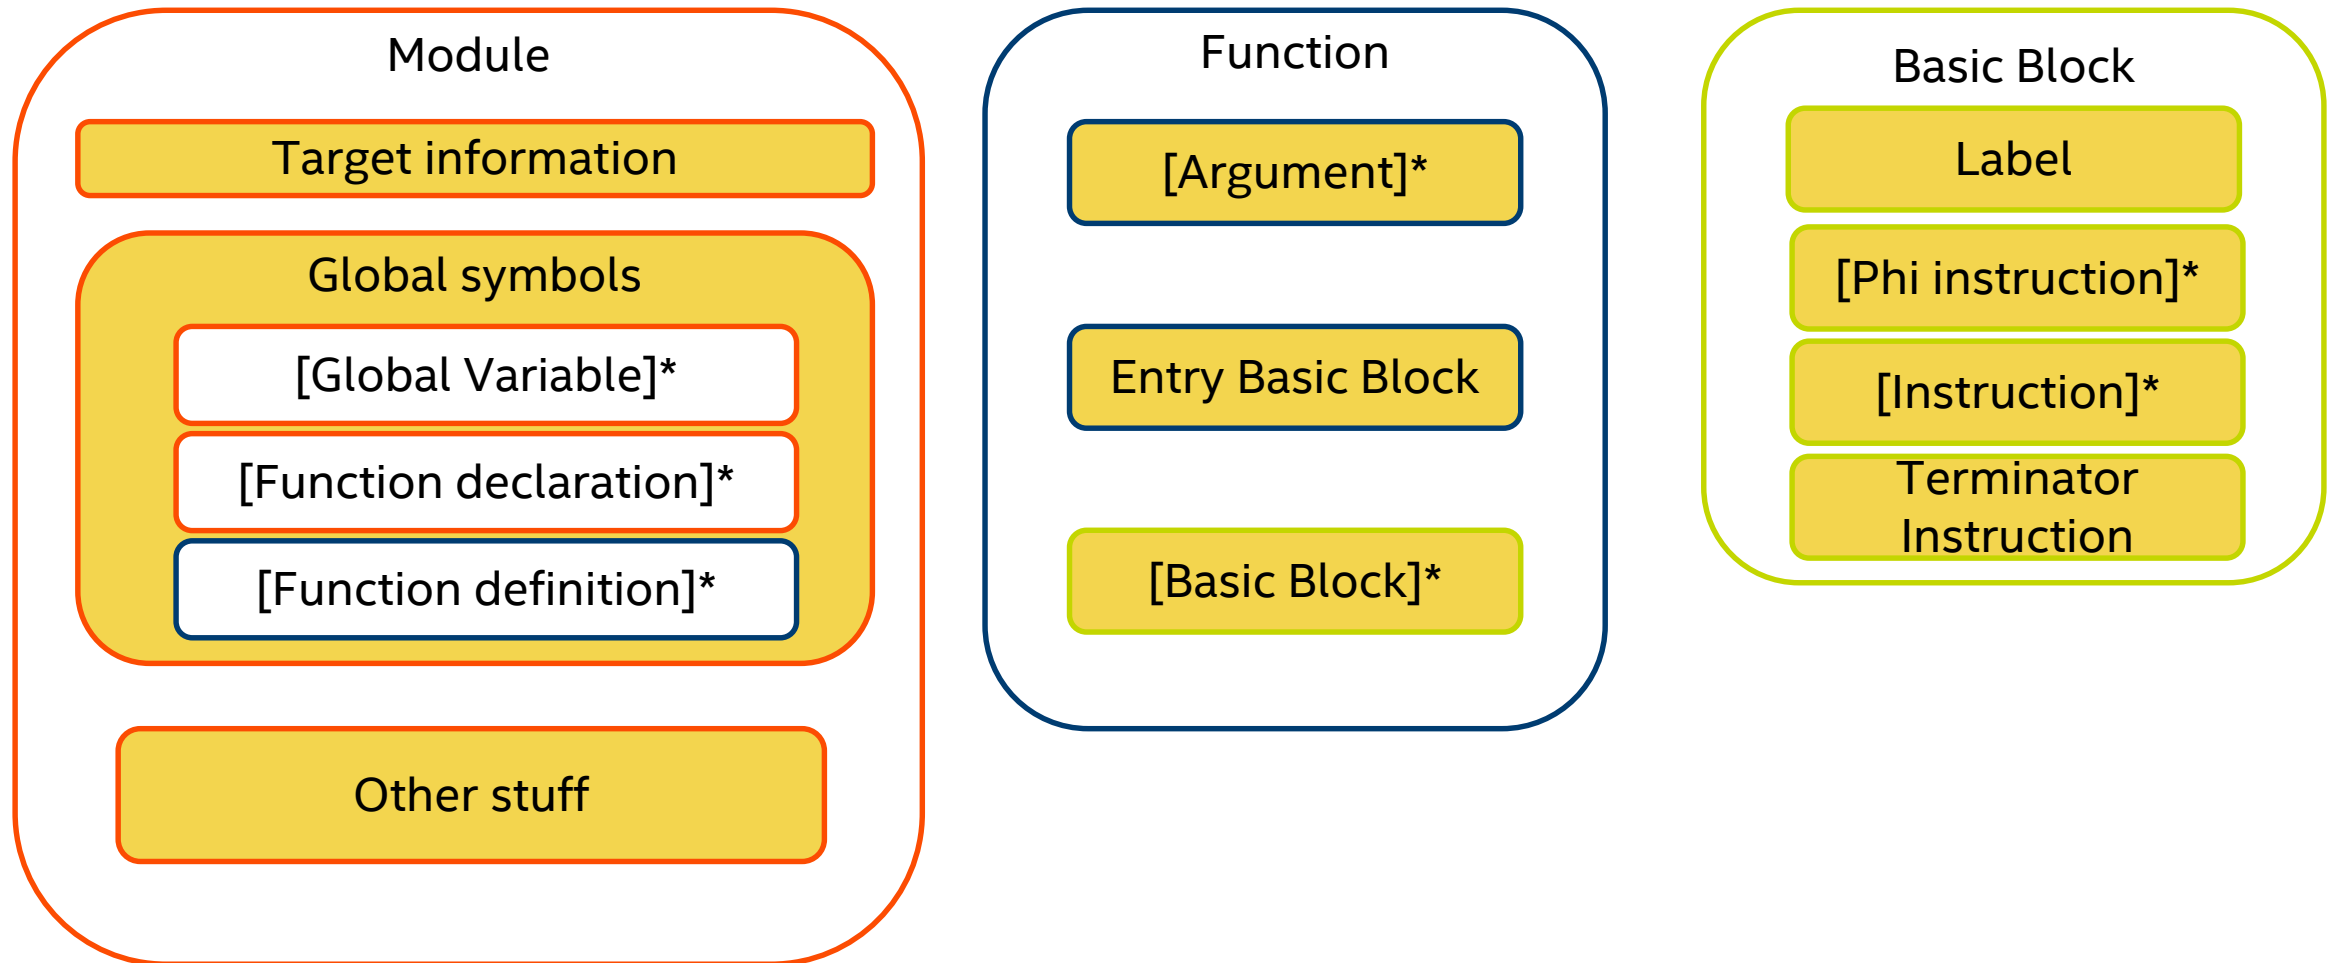
IR representation



IR & the compilation process

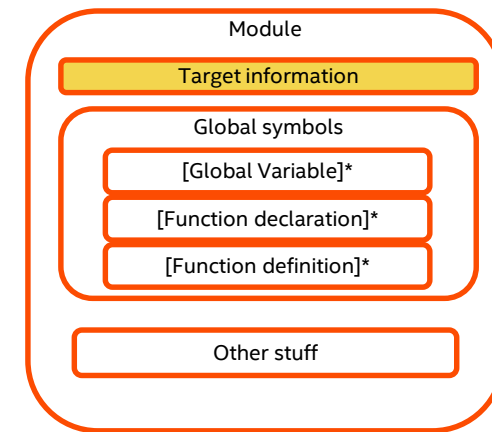


Simplified IR layout



Target information

An IR module usually starts with a pair of strings describing the target:



target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

Diagram annotations:

- Little endian points to 'e' in datalayout.
- ELF mangling points to 'm' in datalayout.
- ABI alignment of i64 points to 'i64:64' in datalayout.
- Native integer widths points to 'n8:16:32:64' in datalayout.
- Architecture points to 'x86_64' in triple.
- Vendor points to 'unknown' in triple.
- System points to 'linux' in triple.
- ABI points to 'gnu' in triple.

A basic main program

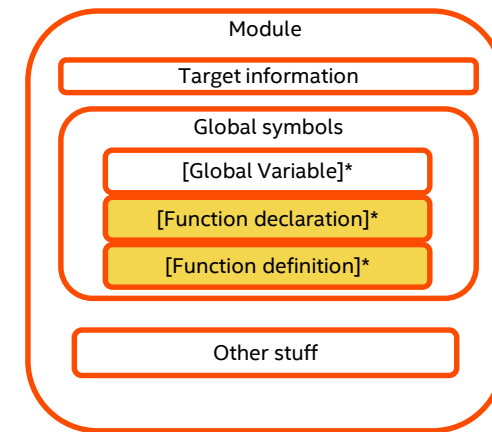
Hand-written IR for this program:

```
int factorial(int val);

int main(int argc, char** argv)
{
    return factorial(2) * 7 == 42;
}
```

declare i32 @factorial(i32)

define i32 @main(i32 %argc, i8** %argv) {
 %1 = call i32 @factorial(i32 2)
 %2 = mul i32 %1, 7
 %3 = icmp eq i32 %2, 42
 %result = zext i1 %3 to i32
 ret i32 %result
}



% Virtual Registers %

Those are “local” variables.

Two flavors of names:

- Unnamed: %<number>
- Named: %<name>

“LLVM IR has infinite registers”

```
declare i32 @factorial(i32)
```

```
define i32 @main(i32 %argc, i8** %argv) {  
    %1 = call i32 @factorial(i32 2)  
    %2 = mul i32 %1, 7  
    %3 = icmp eq i32 %2, 42  
    %result = zext i1 %3 to i32  
    ret i32 %result  
}
```

Types, types everywhere!

Very much a typed language.

```
declare i32 @factorial(i32)

define i32 @main(i32 %argc, i8** %argv) {
    %1 = call i32 @factorial(i32 2)
    %2 = mul i32 %1, 7
    %3 = icmp eq i32 %2, 42
    %result = zext i1 %3 to i32
    ret i32 %result
}
```

Types, types everywhere!

Very much a typed language.

```
    i32                                i32
    i32      i32      i32      i8**      ) {
      i32
    i32
      i32
        i1      i32
i32
}
```

Types, types everywhere!

The instructions explicitly dictate the types expected.

Easy to figure out argument types.

```
declare i32 @factorial(i32)

define i32 @main(i32 %argc, i8** %argv) {
    %1 = call i32 @factorial(i32 2)
    %2 = mul i32 %1, 7
    %3 = icmp eq i32 %2, 42
    %result = zext i1 %3 to i32
    ret i32 %result
}
```

Types, types everywhere!

The instructions explicitly dictate the types expected.

Easy to figure out argument types.

Easy to figure out return types (mostly)

```
declare i32 @factorial(i32)

define i32 @main(i32 %argc, i8** %argv) {
    %1 = call i32 @factorial(i32 2)
    %2 = mul i32 %1, 7
    %3 = icmp eq i32 %2, 42
    %result = zext i1 %3 to i32
    ret i32 %result
}
```

Types, types everywhere!

No implicit conversions!

```
declare i32 @factorial(i32)

define i32 @main(i32 %argc, i8** %argv) {
    %1 = call i32 @factorial(i32 2)
    %2 = mul i32 %1, 7
    %3 = icmp eq i32 %2, 42
    %result = zext i1 %3 to i32
    ret i32 %result
}
```

Types, types everywhere!

No implicit conversions!

To check if this is valid IR:

`opt -verify input.ll`

```
declare i32 @factorial(i32)

define i32 @main(i32 %argc, i8** %argv) {
    %1 = call i32 @factorial(i32 2)
    %2 = mul i32 %1, 7
    %3 = icmp eq i32 %2, 42

    ret i32 %3
}
```

```
opt: test.ll:8:11: error: '%3' defined with type 'i1' but expected 'i32'
```

The LangRef is your friend

Instructions often have many variants.

What else could a call instruction possibly need?

```
declare i32 @factorial(i32)

define i32 @main(i32 %argc, i8** %argv) {
    %1 = call i32 @factorial(i32 2)
    %2 = mul i32 %1, 7
    %3 = icmp eq i32 %2, 42
    %result = zext i1 %3 to i32
    ret i32 %result
}
```


The LangRef is your friend

'call' Instruction

Syntax: `%1 = call i32 @factorial(i32 2)`

```
<result> = [tail | musttail | notail ] call [fast-math flags] [cconv] [ret attrs] [addrspace(<num>)]  
          [<ty>|<fnty> <fnptrval>(<function args>)] [fn attrs] [ operand bundles ]
```

Overview:

The 'call' instruction represents a simple function call.

Arguments:

This instruction requires several arguments:

The LangRef is your friend

Semantics:

The 'call' instruction is used to cause control flow to transfer to a specified function, with its incoming arguments bound to the specified values. Upon a 'ret' instruction in the called function, control flow continues with the instruction after the function call, and the return value of the function is bound to the result argument.

Example:

```
%retval = call i32 @test(i32 %argc)
call i32 (i8*, ...)* @printf(i8* %msg, i32 12, i8 42)      ; yields i32
%X = tail call i32 @foo()                                ; yields i32
%Y = tail call fastcc i32 @foo() ; yields i32
call void %foo(i8 97 signext)

%struct.A = type { i32, i8 }
%r = call %struct.A @foo()                                ; yields { i32, i8 }
%gr = extractvalue %struct.A %r, 0                        ; yields i32
%gr1 = extractvalue %struct.A %r, 1                       ; yields i8
%Z = call void @foo() noreturn                            ; indicates that %foo never returns normally
%ZZ = call zeroext i32 @bar()                             ; Return value is %zero extended
```

Recursive factorial

```
// Precondition: val is non-negative.
int factorial(int val) {
    if (val == 0)
        return 1;
    return val * factorial(val - 1);
}
```

```
; Precondition: %val is non-negative.
define i32 @factorial(i32 %val) {
    %is_base_case = icmp eq i32 %val, 0
    br i1 %is_base_case, label %base_case, label %recursive_case

base_case:
    ret i32 1

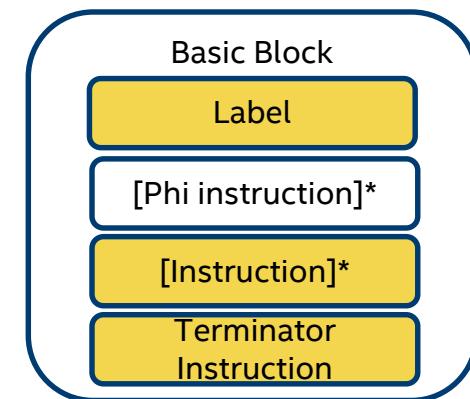
recursive_case:
    %1 = add i32 -1, %val
    %2 = call i32 @factorial(i32 %0)
    %3 = mul i32 %val, %1
    ret i32 %2
}
```

Basic Blocks

List of non-terminator instructions ending with a terminator instruction:

- **Branch - “br”**
- **Return - “ret”**
- **Switch – “switch”**
- **Unreachable – “unreachable”**
- **Exception handling instructions**

```
; Precondition: %val is non-negative.
define i32 @factorial(i32 %val) {
    %is_base_case = icmp eq i32 %val, 0
    br i1 %is_base_case, label %base_case, label %recursive_case
base_case:
    ret i32 1
recursive_case:
    %1 = add i32 -1, %val
    %2 = call i32 @factorial(i32 %0)
    %3 = mul i32 %val, %1
    ret i32 %2
}
```



Basic Blocks


List of non-terminator instructions ending with a terminator instruction:

- Return - “ret”

Execution proceeds to:

- calling function

```
; Precondition: %val is non-negative.
define i32 @factorial(i32 %val) {
    %is_base_case = icmp eq i32 %val, 0
    br i1 %is_base_case, label %base_case, label %recursive_case
base_case:
    ret i32 1
recursive_case:
    %1 = add i32 -1, %val
    %2 = call i32 @factorial(i32 %0)
    %3 = mul i32 %val, %1
    ret i32 %2
}
```



Basic Blocks

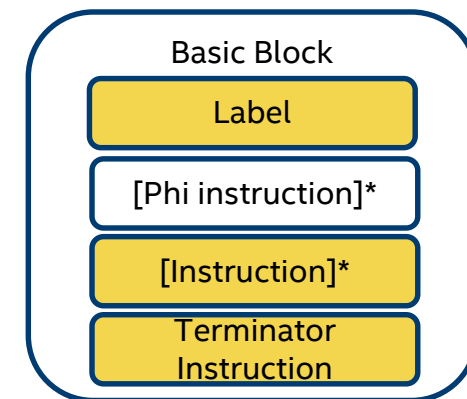
List of non-terminator instructions ending with a terminator instruction:

- Branch - “br”

Execution proceeds to:

- another Basic Block
 - It's successor!

```
; Precondition: %val is non-negative.
define i32 @factorial(i32 %val) {
    %is_base_case = icmp eq i32 %val, 0
    br i1 %is_base_case, label %base_case, label %recursive_case
base_case:
    ret i32 1
recursive_case:
    %1 = add i32 -1, %val
    %2 = call i32 @factorial(i32 %0)
    %3 = mul i32 %val, %1
    ret i32 %2
}
```



Control Flow Graph (CFG)

; Precondition: %val is non-negative.

```
define i32 @factorial(i32 %val) {
```

```
entry:
```

```
    %is_base_case = icmp eq i32 %val, 0
```

```
    br i1 %is_base_case, label %base_case, label %recursive_case
```

```
base_case:
```

```
    ret i32 1
```

```
recursive_case:
```

```
    %0 = add i32 -1, %val
```

```
    %1 = call i32 @factorial(i32 %0)
```

```
    %2 = mul i32 %val, %1
```


```
    ret i32 %2
```

```
}
```

```
; preds = %entry
```

```
; preds = %entry
```

Automatically generated comments




Control Flow Graph (CFG)

The optimizer can generate the CFG in dot format:

```
opt -analyze -dot-cfg-only  
<input.ll>
```

-dot-cfg-only = Generate .dot files.
Don't include instructions.

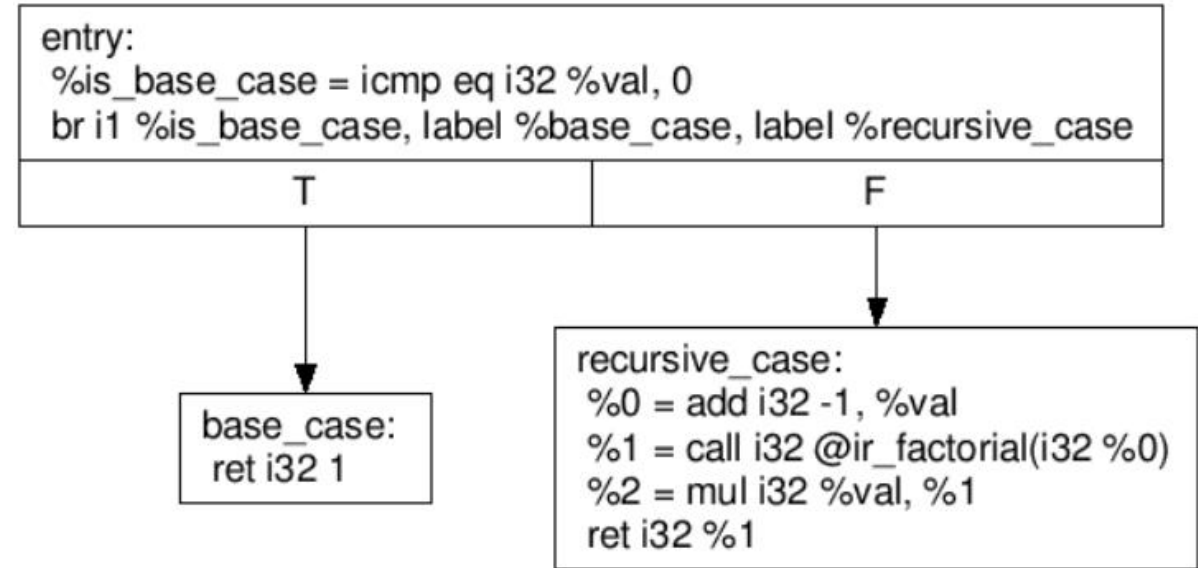


Control Flow Graph (CFG)

The optimizer can generate the CFG in dot format:

`opt -analyze -dot-cfg <input.ll>`

`-dot-cfg` = Generate .dot files.



CFG for 'ir_factorial' function

Implicit labels

Every Basic Block has a label...

... even if it's not explicit

Implicit labels

Every Basic Block has a label...


... even if it's not explicit

```
; Precondition: %val is non-negative.
define i32 @factorial(i32 %val) {
  entry:
    %is_base_case = icmp eq i32 %val, 0
    br i1 %is_base_case, label %base_case, label %recursive_case
  base_case:
    ret i32 1
  recursive_case:
    %0 = add i32 -1, %val
    %1 = call i32 @factorial(i32 %0)
    %2 = mul i32 %val, %1
    ret i32 %2
}
```

Implicit labels

Every Basic Block has a label...
... even if it's not explicit

```
opt: ir_implementation.ll:11:3: error: instruction expected to be numbered '%1'
    %0 = add i32 -1, %val
    ^
; Precondition: %val is non-negative.
define i32 @factorial(i32 %val) {
    %0: ; Implicit label!
    %is_base_case = icmp eq i32 %val, 0
    ; ...
recursive_case:
    %0 = add i32 -1, %val
    %1 = call i32 @factorial(i32 %0)
    %2 = mul i32 %val, %1
    ret i32 %2
}
```



Implicit labels

Every Basic Block has a label...

... even if it's not explicit

```
opt: ir_implementation.ll:11:3: error: instruction expected to be numbered '%1'
    %0 = add i32 -1, %val
    ^
```

; Precondition: %val is non-negative.

```
define i32 @factorial(i32 %val) {
```

```
    %0: ; Implicit label!
```

```
    %is_base_case = icmp eq i32 %val, 0
```

```
    recursive_case:
```

```
        %1 = add i32 -1, %val
```

```
        %2 = call i32 @factorial(i32 %1)
```

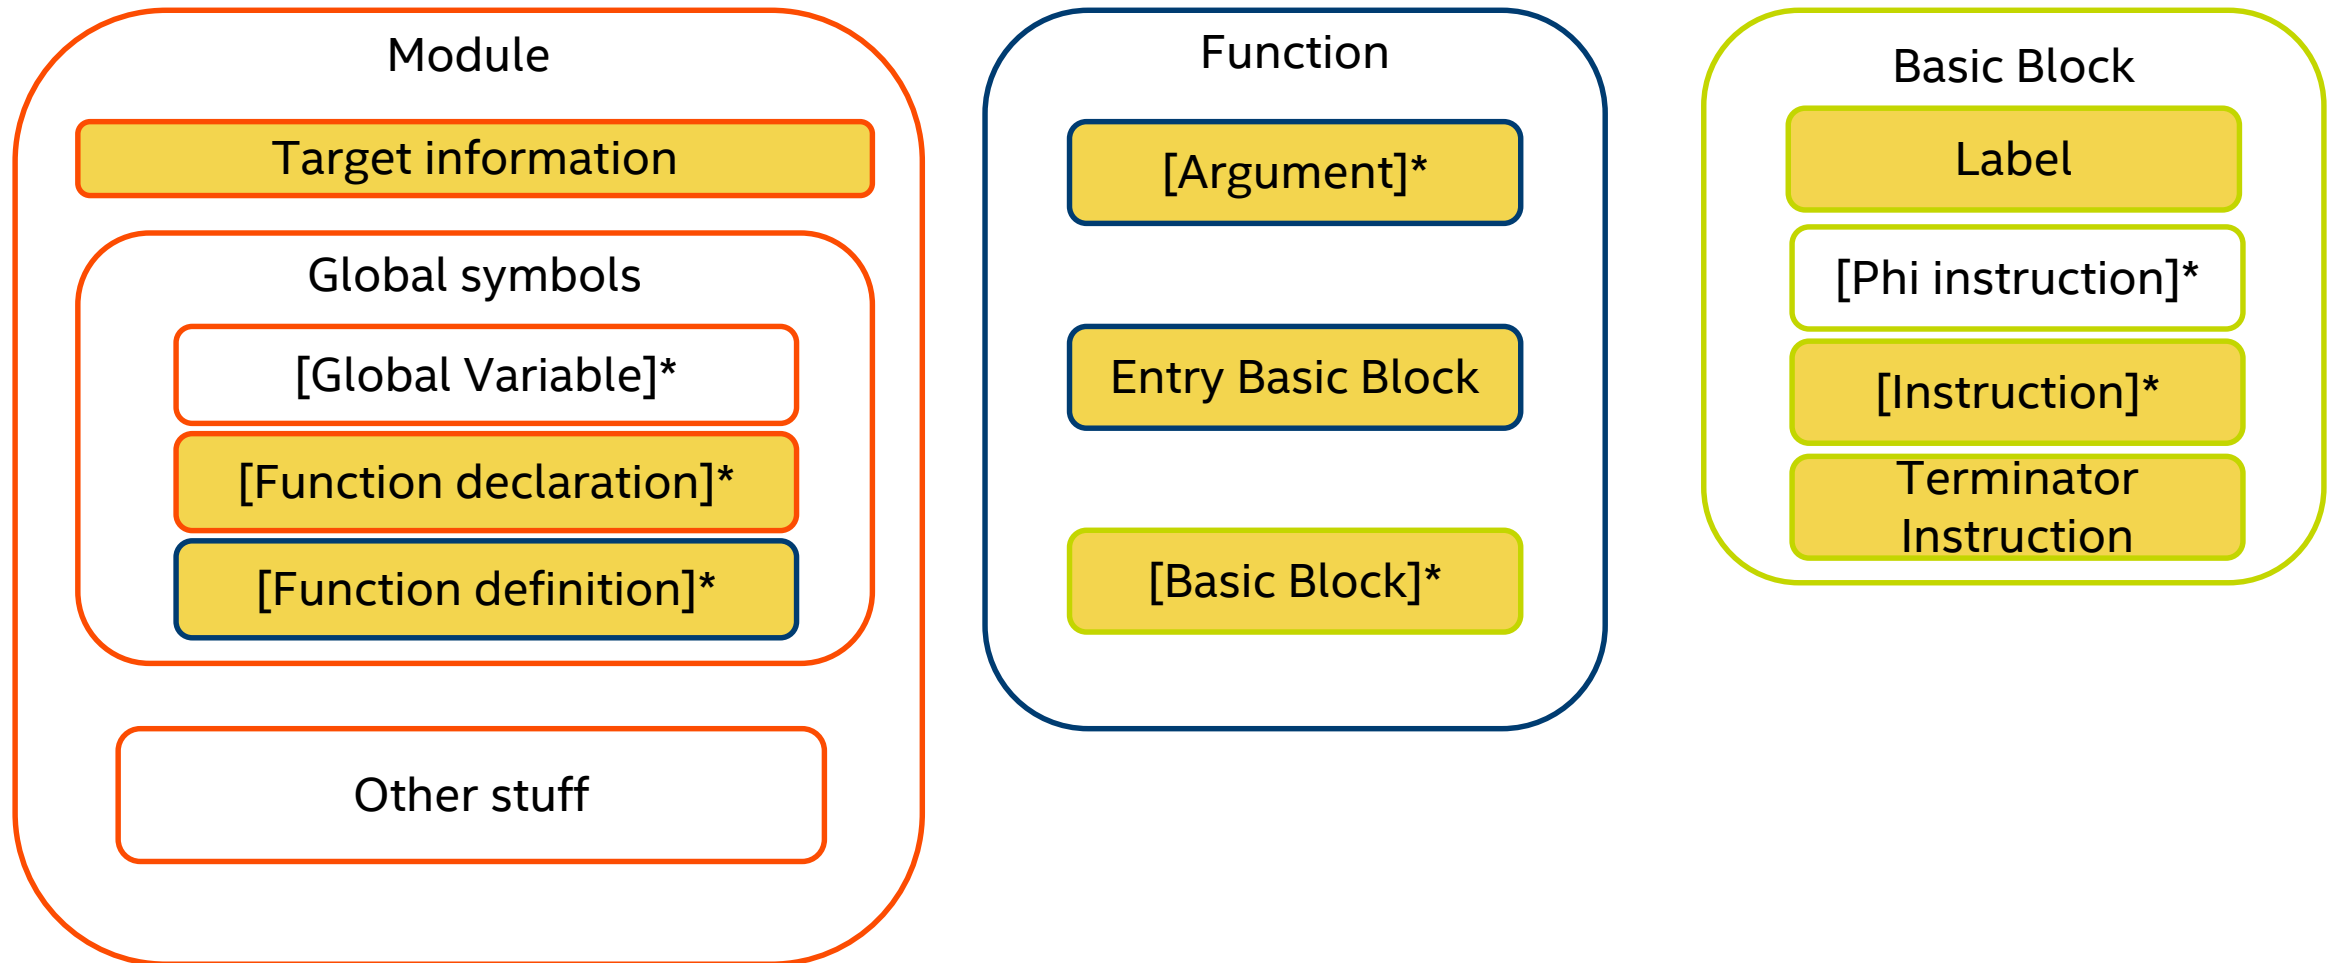
```
        %3 = mul i32 %val, %2
```

```
        ret i32 %3
```

```
}
```

And the same is true for function arguments!

Simplified IR layout



Iterative factorial

```
int factorial(int val) {  
    int temp = 1;  
    for (int i = 2; i <= val; ++i)  
        temp *= i;  
    return temp;  
}
```

You wish you could do this...

```
define i32 @factorial(i32 %val) {  
entry:  
    %i = add i32 0, 2  
    %temp = add i32 0, 1  
    br label %check_for_condition  
check_for_condition:  
  
    %i_leq_val = icmp sle i32 %i, %val  
    br i1 %i_leq_val, label %for_body, label %end_loop  
for_body:  
  
    %temp = mul i32 %temp, %i  
    %i = add i32 %i, 1  
    br label %check_for_condition  
end_loop:  
    ret i32 %temp  
}
```

Iterative factorial

```
int factorial(int val) {  
    int temp = 1;  
    for (int i = 2; i <= val; ++i)  
        temp *= i;  
    return temp;  
}
```

```
define i32 @factorial(i32 %val) {  
entry:  
    %i = add i32 0, 2  
    %temp = add i32 0, 1  
    br label %check_for_condition  
check_for_condition:  
  
    %i_leq_val = icmp sle i32 %i, %val  
    br i1 %i_leq_val, label %for_body, label %end_loop  
for_body:
```

You wish you could do this...

```
{ %temp = mul i32 %temp, %i  
  %i = add i32 %i, 1
```



```
opt: test.ll:12:5: error: multiple definition of local value named 'temp'  
    %temp = mul i32 %temp, %i  
    ^
```

```
}
```


Static Single Assignment (SSA)

Every variable is assigned *exactly* once.

Every variable is defined before it is used.

Iterative factorial

```
int factorial(int val) {  
    int temp = 1;  
    for (int i = 2; i <= val; ++i)  
        temp *= i;  
    return temp;  
}
```

```
define i32 @factorial(i32 %val) {  
entry:
```

```
    %i = add i32 0, 2
```

```
    %temp = add i32 0, 1
```

```
    br label %check_for_condition
```

```
check_for_condition:
```

```
    %i_leq_val = icmp sle i32 %i, %val
```

```
    br i1 %i_leq_val, label %for_body, label %end_loop
```

```
for_body:
```

```
    %new_temp = mul i32 %temp, %i
```

```
    %i_plus_one = add i32 %i, 1
```

```
    br label %check_for_condition
```

```
end_loop:
```

```
    ret i32 %temp
```

```
}
```


Now %i is always 2!

So you do this:



Now %temp is always 1!

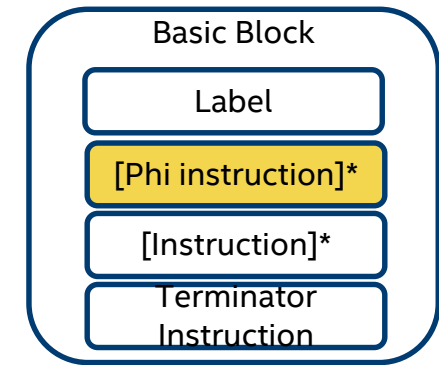
Phi to the rescue!



`<result> = phi <ty> [, [] ...`

Select a value based on the **BasicBlock** that executed previously!


Phi to the rescue!



`<result> = phi <ty> [<val0>, <label0>], [` `]` ...

Select a value based on the **BasicBlock** that executed previously!

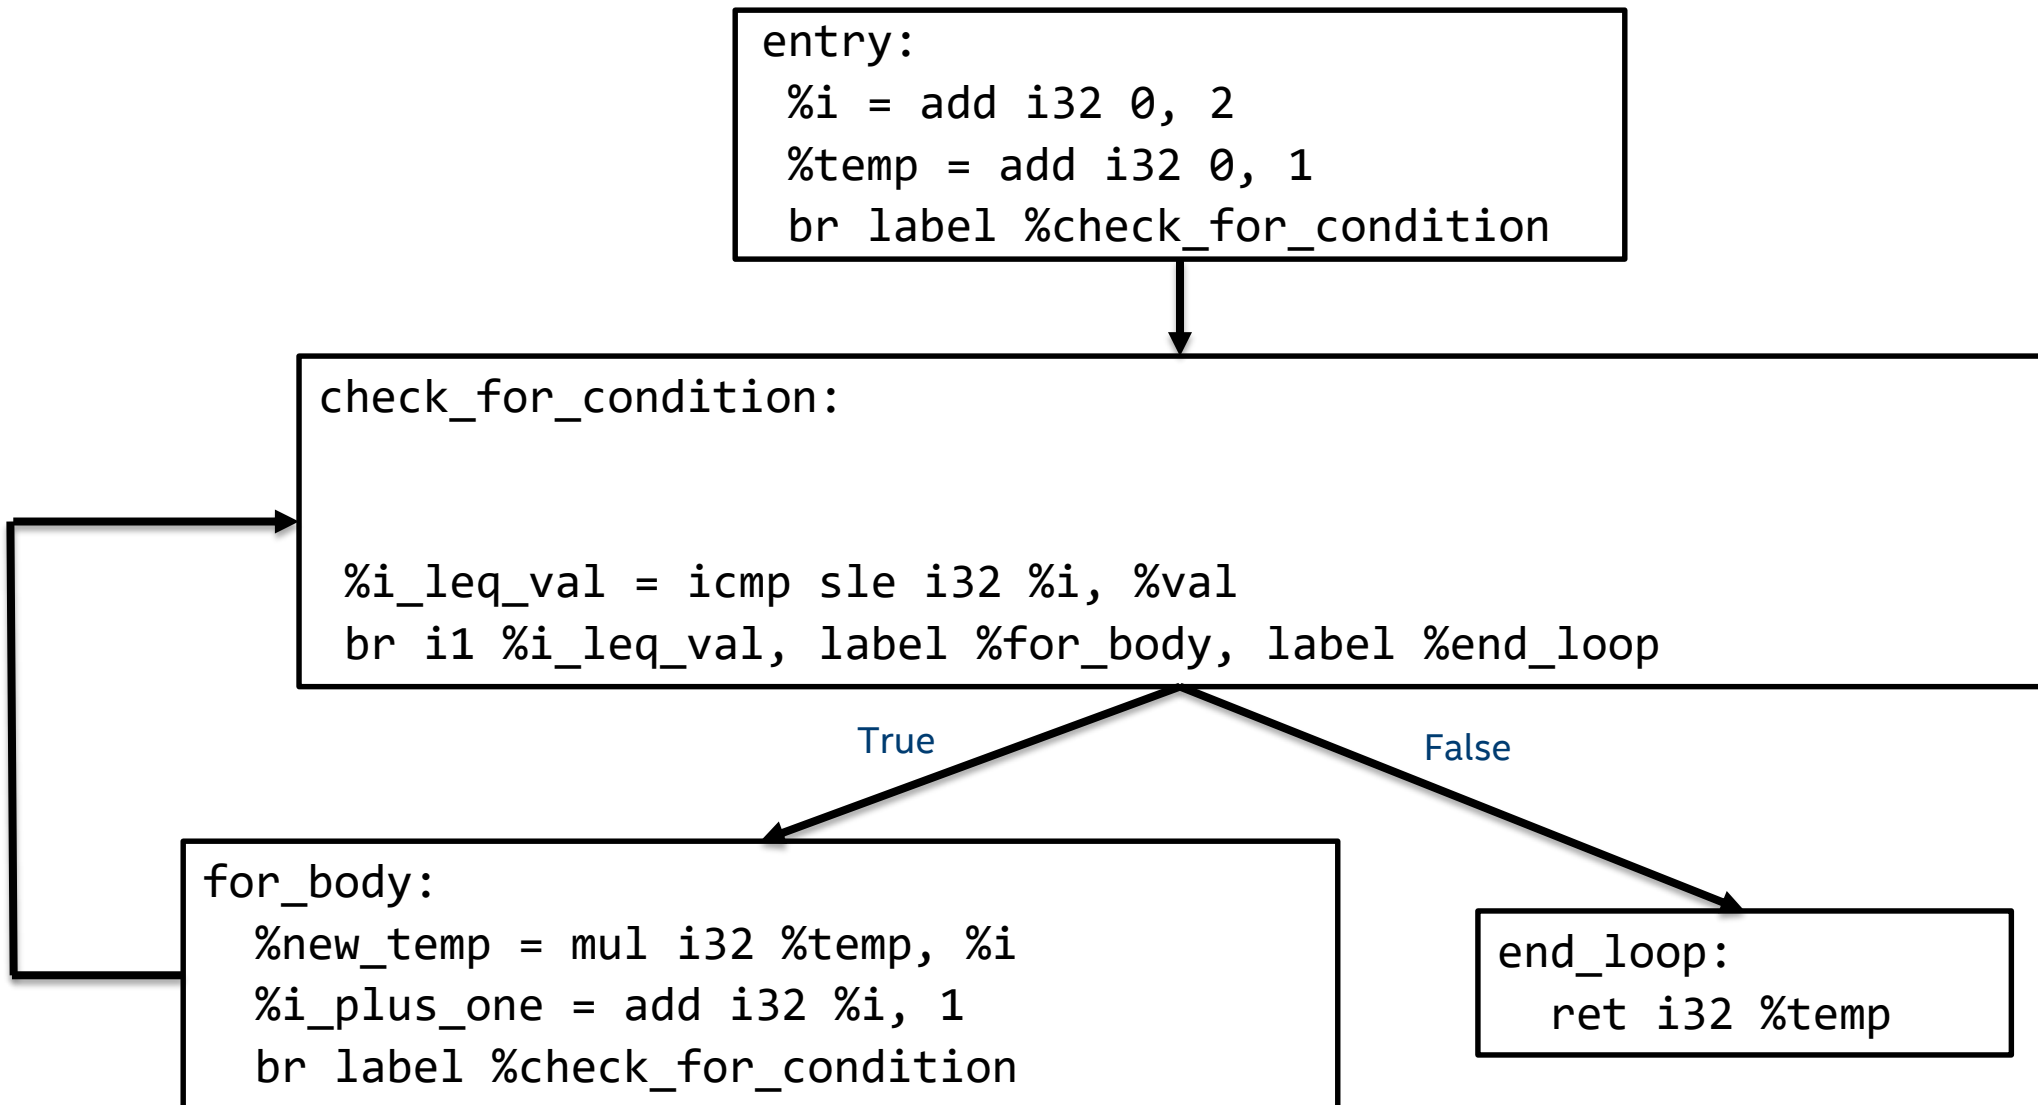
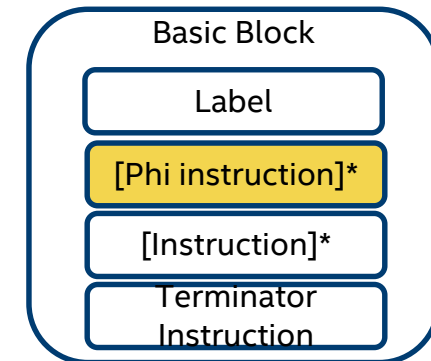
Phi to the rescue!



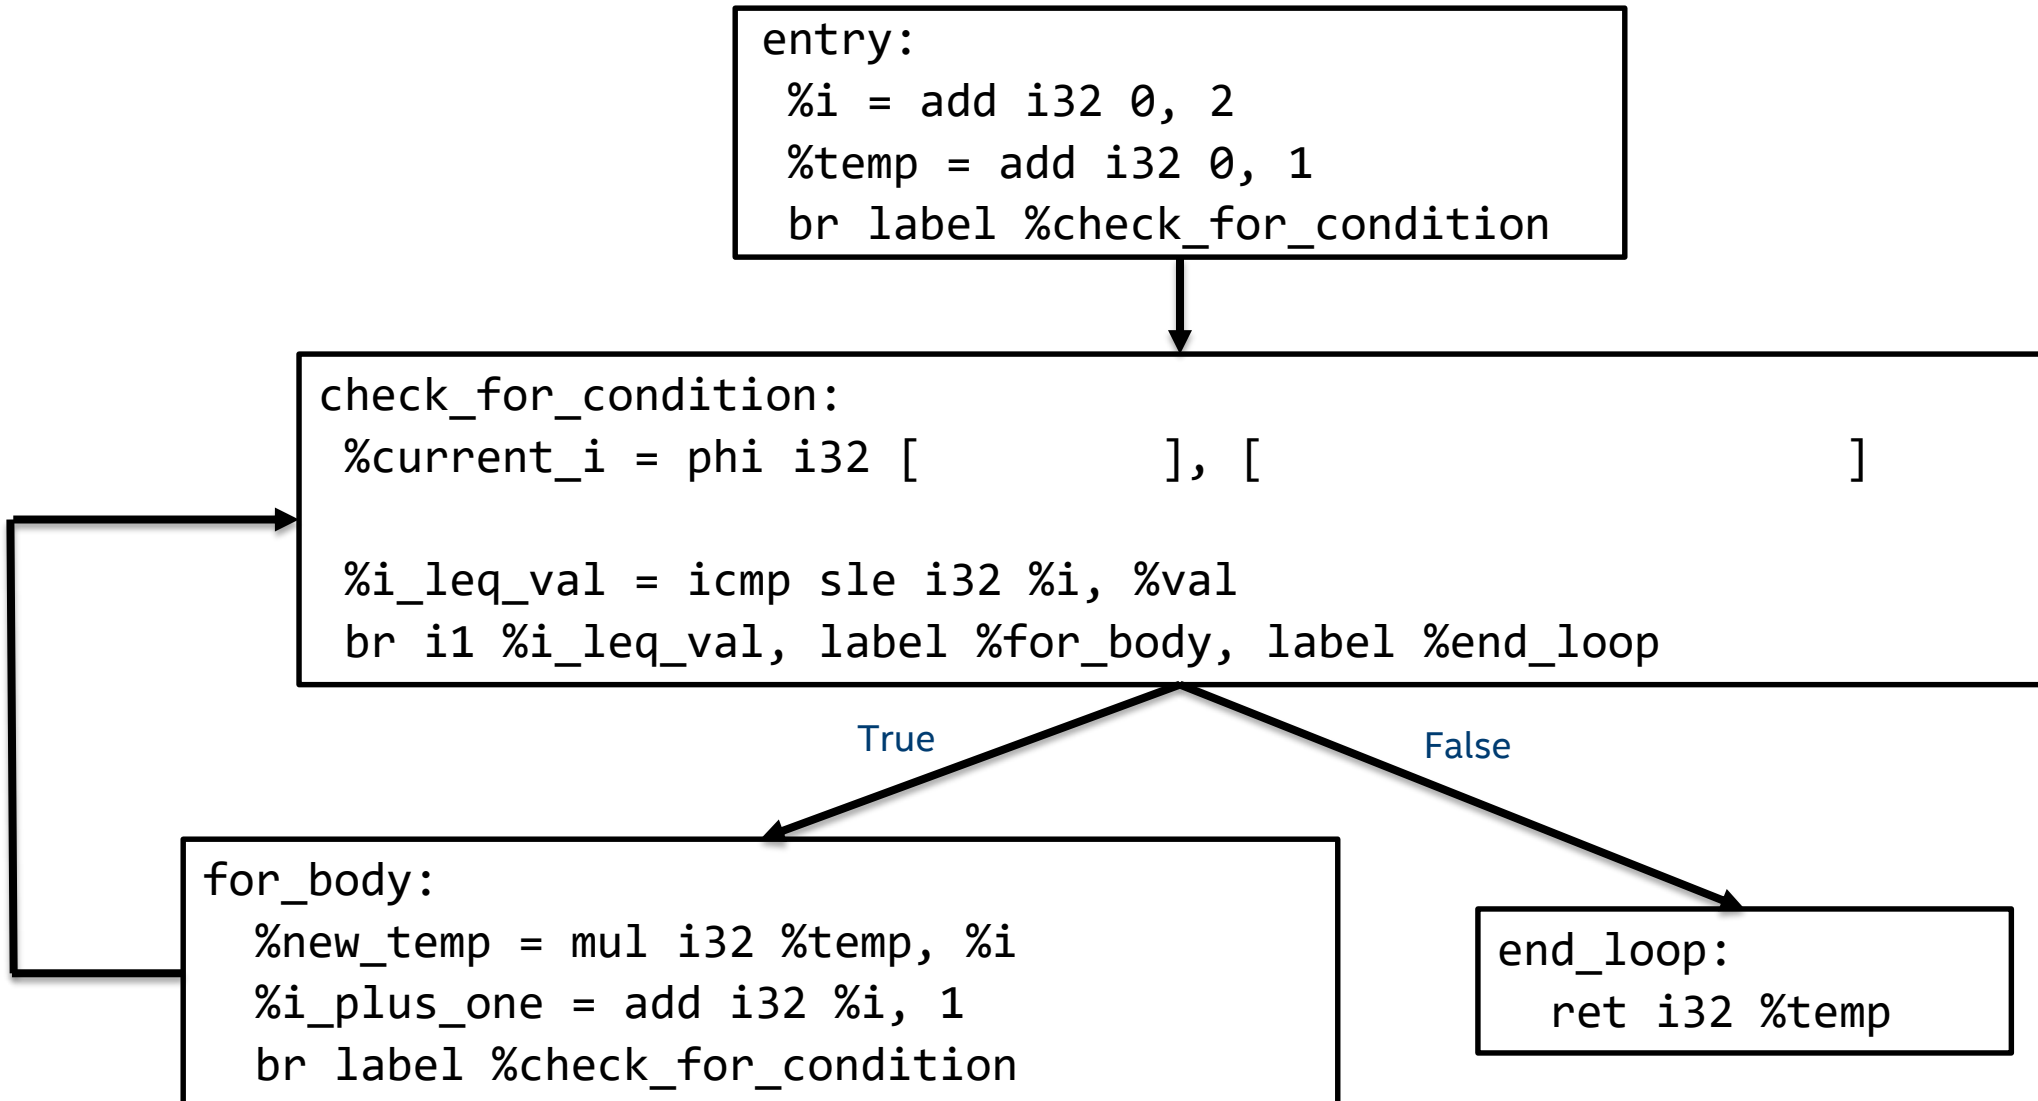

`<result> = phi <ty> [<val0>, <label0>], [<val1>, <label1>] ...`

Select a value based on the **BasicBlock** that executed previously!


Phis to the rescue!



Phis to the rescue!



Phis to the rescue!



```
entry:  
  %i = add i32 0, 2  
  %temp = add i32 0, 1  
  br label %check_for_condition
```

```
check_for_condition:  
  %current_i = phi i32 [2, %entry], [  
    ]  
  
  %i_leq_val = icmp sle i32 %i, %val  
  br i1 %i_leq_val, label %for_body, label %end_loop
```

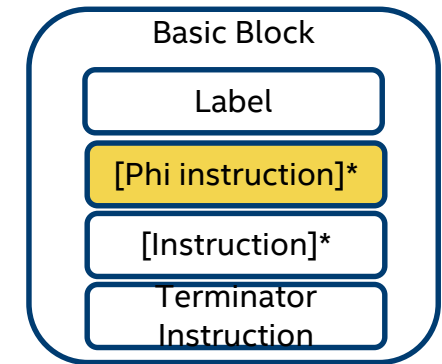
True

False

```
for_body:  
  %new_temp = mul i32 %temp, %i  
  %i_plus_one = add i32 %i, 1  
  br label %check_for_condition
```

```
end_loop:  
  ret i32 %temp
```


Phis to the rescue!



```
entry:  
  %i = add i32 0, 2  
  %temp = add i32 0, 1  
  br label %check_for_condition
```

```
check_for_condition:  
  %current_i = phi i32 [2, %entry], [%i_plus_one, %for_body]  
  
  %i_leq_val = icmp sle i32 %i, %val  
  br i1 %i_leq_val, label %for_body, label %end_loop
```

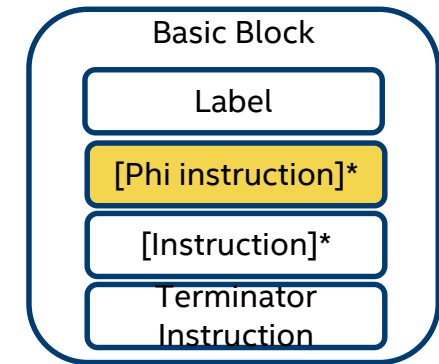
True

False

```
for_body:  
  %new_temp = mul i32 %temp, %i  
  %i_plus_one = add i32 %i, 1  
  br label %check_for_condition
```

```
end_loop:  
  ret i32 %temp
```

Phis to the rescue!



```
entry:  
  %i = add i32 0, 2  
  %temp = add i32 0, 1  
  br label %check_for_condition
```

```
check_for_condition:  
  %current_i = phi i32 [2, %entry], [%i_plus_one, %for_body]  
  
  %i_leq_val = icmp sle i32 %current_i, %val  
  br i1 %i_leq_val, label %for_body, label %end_loop
```

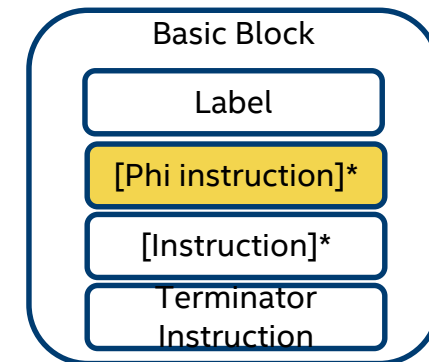
True

False

```
for_body:  
  %new_temp = mul i32 %temp, %current_i  
  %i_plus_one = add i32 %current_i, 1  
  br label %check_for_condition
```

```
end_loop:  
  ret i32 %temp
```

Phis to the rescue!



entry:

```
%temp = add i32 0, 1  
br label %check_for_condition
```

check_for_condition:

```
%current_i = phi i32 [2, %entry], [%i_plus_one, %for_body]
```

```
%i_leq_val = icmp sle i32 %current_i, %val  
br i1 %i_leq_val, label %for_body, label %end_loop
```

True

False

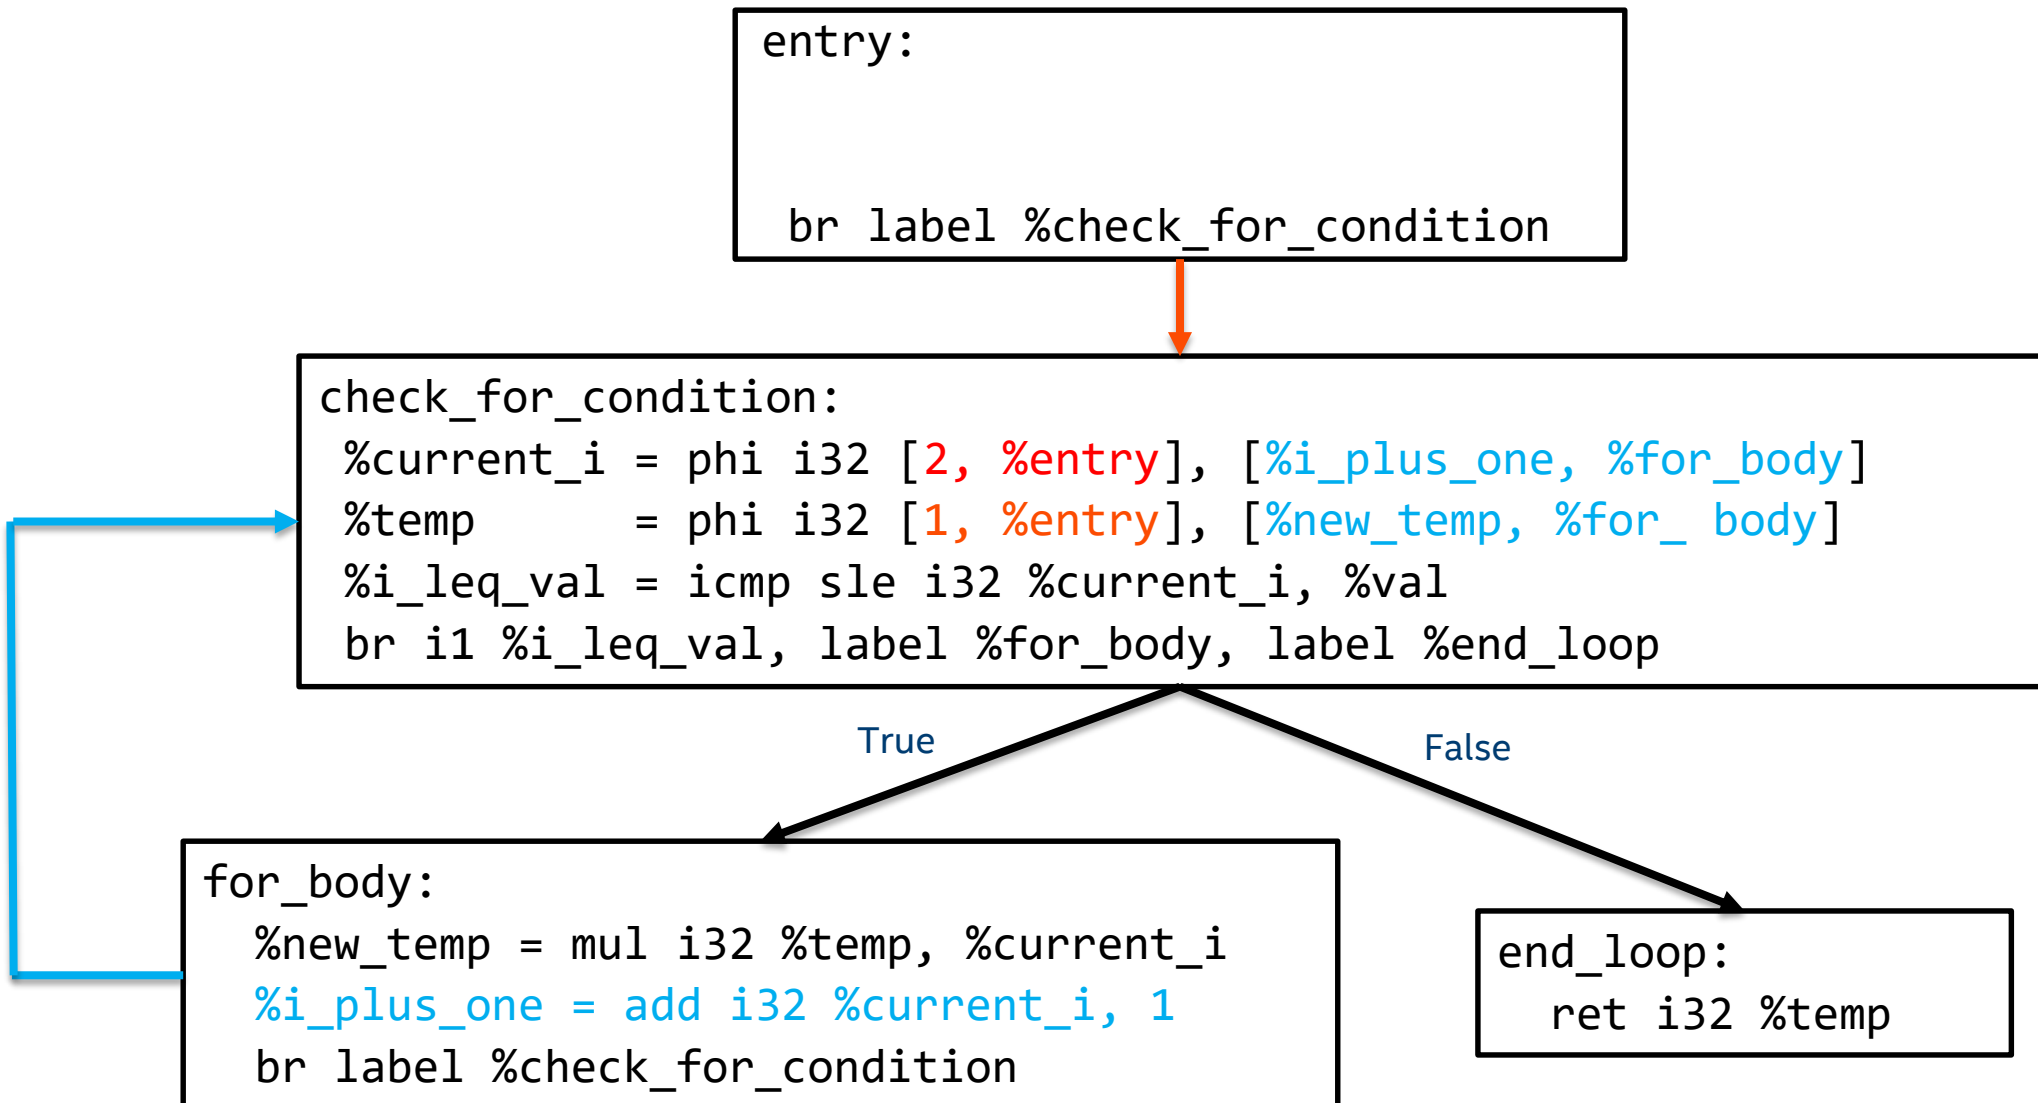
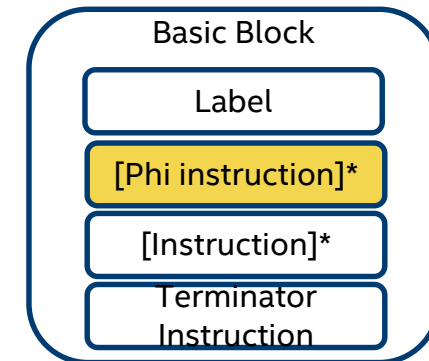
for_body:

```
%new_temp = mul i32 %temp, %current_i  
%i_plus_one = add i32 %current_i, 1  
br label %check_for_condition
```

end_loop:

```
ret i32 %temp
```

Phis to the rescue!



Another way to cheat SSA

Frontend generates something different!

Gets around SSA restriction by writing to memory.

<https://godbolt.org/z/Nlx6T5>

(remember to untick the “hide comments” option!)

But the optimized code is similar to what we had (with O1):

<https://godbolt.org/z/OirW9y>

Another way to cheat SSA

Alloca instruction:

- You give it a type, it gives you a pointer to that type:
 - `%ptr = alloca i32` ; ptr is i32*
 - `%ptr = alloca <any_type>` ; ptr is <any_type>*
- Allocates memory on the stack frame of the executing function.
- Automatically released.
 - Akin to changing the stack pointer.
- Plays a big part in generating IR in SSA form.

Allocas to the rescue!

```
entry:  
  %i.addr = alloca i32  
  %temp.addr = alloca i32  
  store i32 2, i32* %i.addr  
  store i32 1, i32* %temp.addr  
  br label %check_for_condition
```

```
check_for_condition:  
  
  %i_leq_val = icmp sle i32 %current_i, %val  
  br i1 %i_leq_val, label %for_body, label %end_loop
```

True

False

```
for_body:  
  %i_plus_one = add i32 %current_i, 1  
  %new_temp = mul i32 %temp, %current_i  
  store i32 %i_plus_one, i32* %i.addr  
  store i32 %new_temp, i32* %temp.addr  
  br label %check_for_condition
```

```
end_loop:  
  ret i32 %temp
```

Allocas to the rescue!

```
entry:  
  %i.addr = alloca i32  
  %temp.addr = alloca i32  
  store i32 2, i32* %i.addr  
  store i32 1, i32* %temp.addr  
  br label %check_for_condition
```

```
check_for_condition:  
  %current_i = load i32, i32* %i.addr  
  %temp      = load i32, i32* %temp.addr  
  %i_leq_val = icmp sle i32 %current_i, %val  
  br i1 %i_leq_val, label %for_body, label %end_loop
```

True

False

```
for_body:  
  %i_plus_one = add i32 %current_i, 1  
  %new_temp = mul i32 %temp, %current_i  
  store i32 %i_plus_one, i32* %i.addr  
  store i32 %new_temp, i32* %temp.addr  
  br label %check_for_condition
```

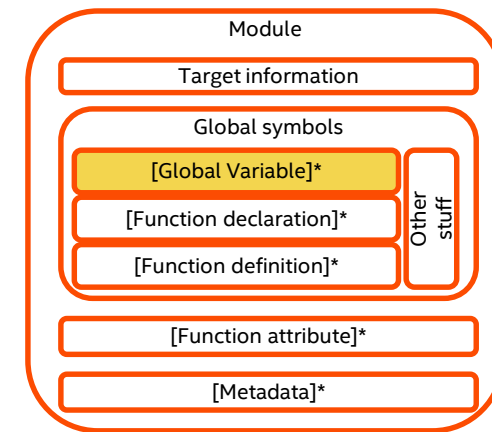
```
end_loop:  
  ret i32 %temp
```


Global variables

Allocas allocate memory for function scopes.

Global variables fill that role for the module in a static way.

- They are always pointers, like the values returned by Allocas.



Global variables

- Name prefixed with “@”.
- Must have a type.
- Must be initialized
- Have the **global** keyword...
- ... xor **constant** (never stored to!)
 - Not to be confused with C++ const

@gv =

@gv = **i8**

@gv = **i8** 42 ; Declarations excepted.

@gv = **global i8** 42

@gv = **constant i8** 42

Global variables

Are always pointers

Always **constant pointers!**

```
@gv = global i16 46  
; ...  
; ... Inside some function:  
%load = load i16 , i16* @gv  
store i16 0, i16* @gv
```

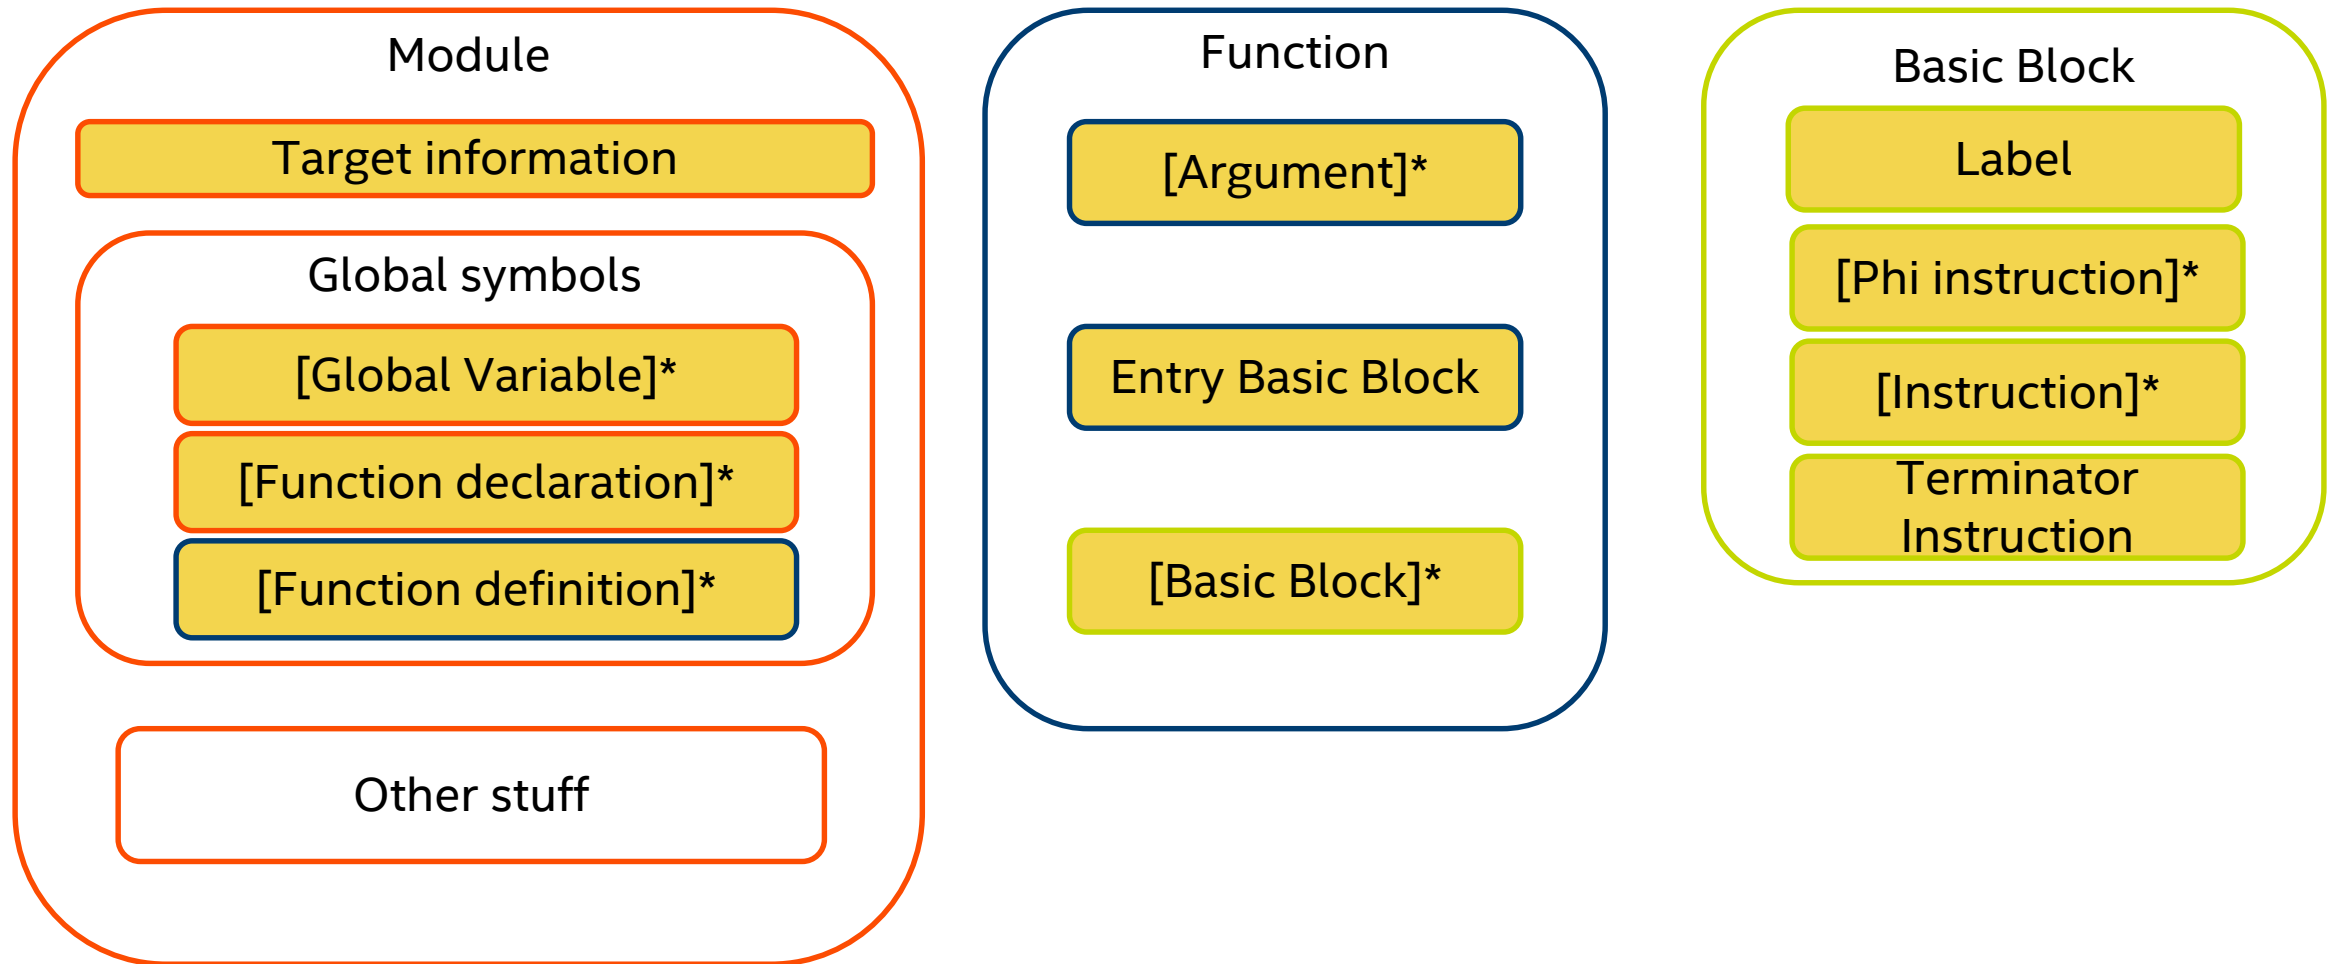
Global variables

Love qualifiers:

```
@<GlobalVarName> = [Linkage] [PreemptionSpecifier] [Visibility]  
                  [DLLStorageClass] [ThreadLocal]  
                  [(unnamed_addr|local_unnamed_addr)] [AddrSpace]  
                  [ExternallyInitialized]  
                  <global | constant> <Type> [<InitializerConstant>]  
                  [, section "name"] [, comdat [($name)]]  
                  [, align <Alignment>] (, !name !N)*
```

Check the language ref.

Simplified IR layout



TYPE SYSTEM AND GEPS!

LLVM's type system

From the language reference:

- Void Type
- Function Type
- First Class Types
 - Single Value Types
 - Integer Type
 - Floating-Point Types
 - X86_mmx Type
 - Pointer Type
 - Vector Type
 - Label Type
 - Token Type
 - Metadata Type
 - Aggregate Types
 - Array Type
 - Structure Type
 - Opaque Structure Types

Aggregate types: arrays

Defined by:

- A constant size. `@array = global [17 x]`
- An element type. `@array = global [17 x i8]`
- [for GVs] an initializer `@array = global [17 x i8] zeroinitializer`

Accessing arrays & manipulating pointers

The Get Element Pointer (GEP) instruction:

- Provides a way to calculate pointer offsets.
- Abstracts away details like:
 - Size of types
 - Padding inside structs
- Intuitive to use...
... once you understand a few basic principles.

Manipulating pointers

The Get Element Pointer (GEP) instruction:


`<result> = getelementptr <ty>, <ty>* <ptrval>, [i32 <idx>]+`

↓
**Base type used for
the first index**

↓
**Base address to start
from**

↓
**Offsets - one per
“dimension”**


Manipulating pointers



`%new_ptr = getelementptr i32, i32* %base, i32 0`

↓
“Offset by 0 elements of the base type”

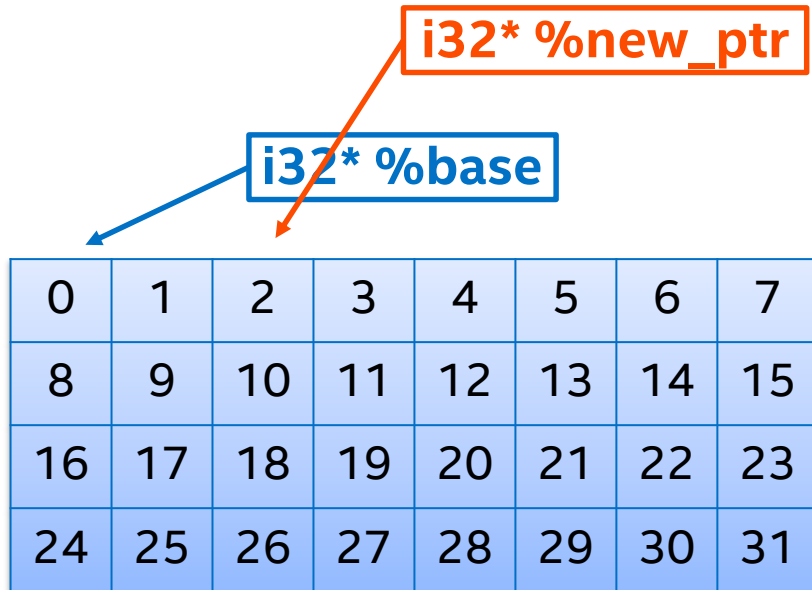
Manipulating pointers



```
%new_ptr = getelementptr i32, i32* %base, i32 1
```

“Offset by 1 elements of the base type”


Manipulating pointers



`%new_ptr = getelementptr i32, i32* %base, i32 2`

“Offset by 2 elements of the base type”


Manipulating pointers



`%new_ptr = getelementptr i32, i32* %base, i32 3`

↓
“Offset by 3 elements of the base type”

Manipulating pointers




`%new_ptr = getelementptr i32, i32* %base, i32 4`

“Offset by 4 elements of the base type”

GEP fundamentals

1. Understand the first index:
 - It does NOT change the resulting pointer type.
 - It offsets by the **base type**.

Manipulating pointers




```
@a_gv = global [6 x i8] zeroinitializer
```

```
%new_ptr = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0
```

“Offset by 0 elements of the base type”

Manipulating pointers




```
@a_gv = global [6 x i8] zeroinitializer
```

```
%new_ptr = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0
```

“Offset by 0 elements of the base type”

Manipulating pointers



```
@a_gv = global [6 x i8] zeroinitializer
```

```
%new_ptr = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 1
```

“Offset by 1 elements of the base type”

GEP fundamentals


1. Understand the first index:

- It does NOT change the pointer type.
- It offsets by the **pointee type**.

2. Further indices:

- Offset inside **aggregate types**. (and vectors)
- Change the pointer type by removing one layer of “aggregation”.

Manipulating pointers



```
@a_gv = global [6 x i8] zeroinitializer
```


```
%new_ptr = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0
```

```
%elem_pt = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0, i32 0
```

“Offset by 0 elements of the base type”

Get the 0th element
from the current
aggregate:
[6 x i8]

Manipulating pointers



```
@a_gv = global [6 x i8] zeroinitializer
```


```
%new_ptr = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0
```

```
%elem_pt = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0, i32 1
```

“Offset by 0 elements of the base type”

Get the 1st element
from the current
aggregate:
[6 x i8]

Manipulating pointers



```
@a_gv = global [6 x i8] zeroinitializer
```

```
%new_ptr = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0
```

```
%elem_pt = getelementptr [6 x i8], [6 x i8]* @a_gv, i32 0, i32 2
```

“Offset by 0 elements of the base type”


Get the 2nd element
from the current
aggregate:
[6 x i8]

Aggregate types: structs

Defined by:

- A name: `%MyStruct =`
- Keyword “type”: `%MyStruct = type`
- A list of types: `%MyStruct = type { i8, i32, [3 x i32] }`

GEPs with structs

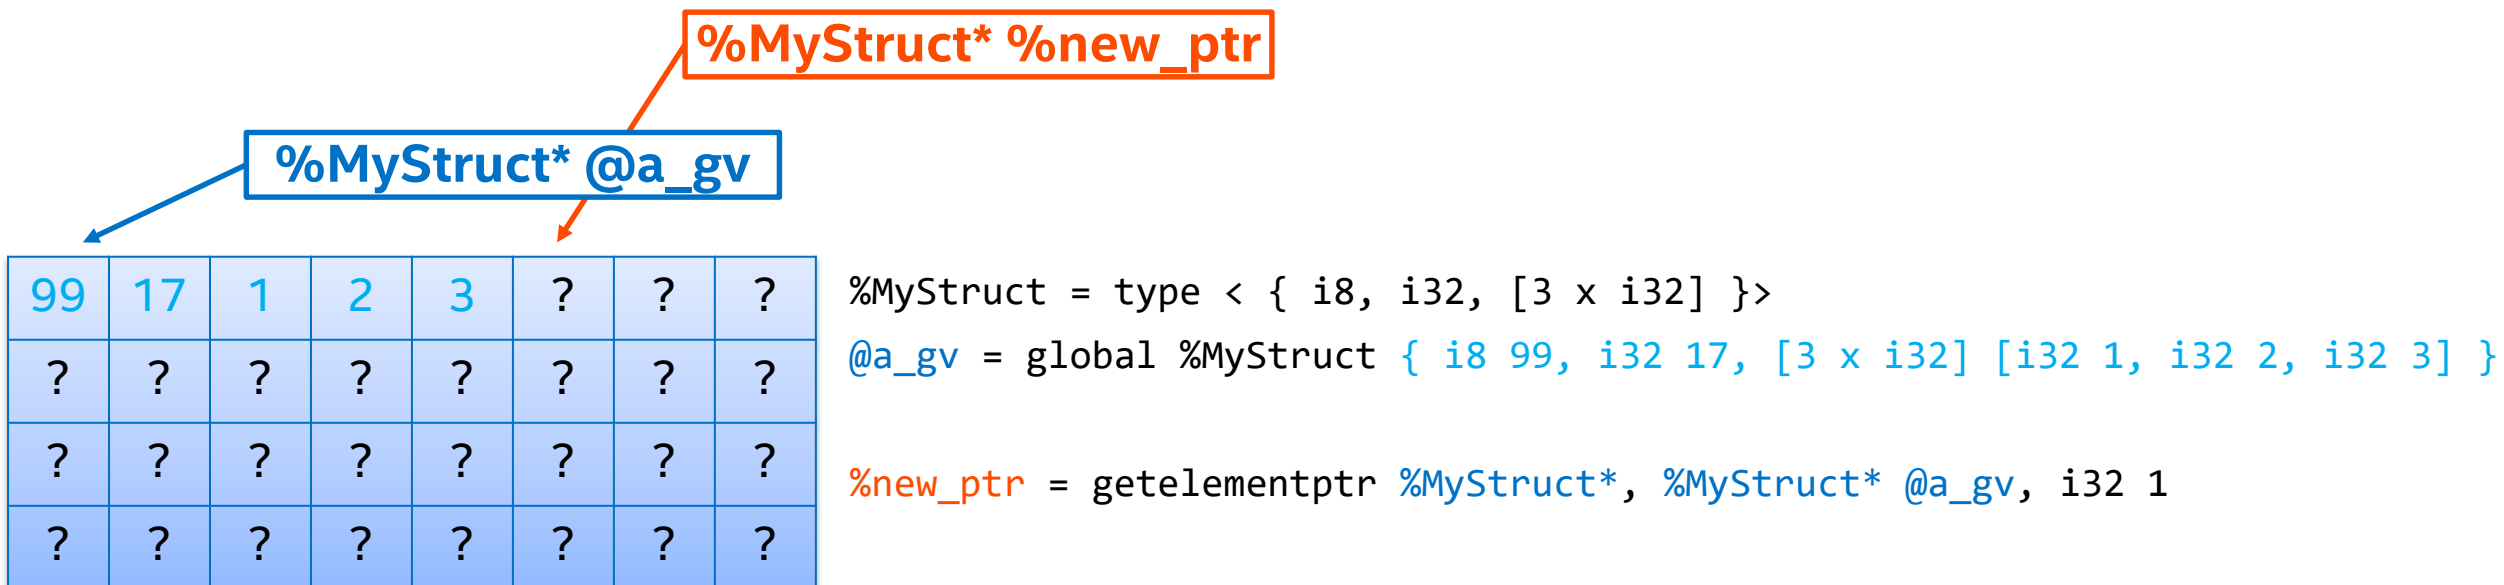


```
%MyStruct = type < { i8, i32, [3 x i32] }>
```


```
@a_gv = global %MyStruct { i8 99, i32 17, [3 x i32] [i32 1, i32 2, i32 3] }
```

```
%new_ptr = getelementptr %MyStruct*, %MyStruct* @a_gv, i32 0
```

GEPs with structs



GEPs with structs



```
%MyStruct = type < { i8, i32, [3 x i32] }>
```

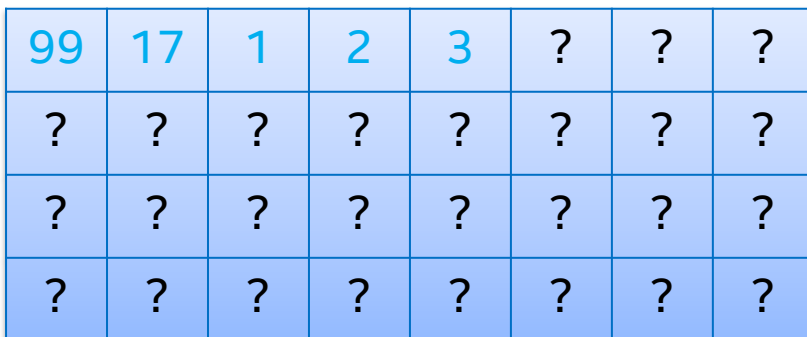
```
@a_gv = global %MyStruct { i8 99, i32 17, [3 x i32] [i32 1, i32 2, i32 3] }
```

```
%new_ptr = getelementptr %MyStruct*, %MyStruct* @a_gv, i32 0
```

GEPs with structs

`i8* %new_ptr`

`%MyStruct* @a_gv`




99	17	1	2	3	?	?	?
?	?	?	?	?	?	?	?
?	?	?	?	?	?	?	?
?	?	?	?	?	?	?	?

```
%MyStruct = type < { i8, i32, [3 x i32] }>
```

```
@a_gv = global %MyStruct { i8 99, i32 17, [3 x i32] [i32 1, i32 2, i32 3] }
```

```
%new_ptr = getelementptr %MyStruct*, %MyStruct* @a_gv, i32 0, i32 0
```

GEPs with structs




```
%MyStruct = type < { i8, i32, [3 x i32] }>
```

```
@a_gv = global %MyStruct { i8 99, i32 17, [3 x i32] [i32 1, i32 2, i32 3] }
```

```
%new_ptr = getelementptr %MyStruct*, %MyStruct* @a_gv, i32 0, i32 1
```

GEPs with structs




```
%MyStruct = type < { i8, i32, [3 x i32] }>
```

```
@a_gv = global %MyStruct { i8 99, i32 17, [3 x i32] [i32 1, i32 2, i32 3] }
```

```
%new_ptr = getelementptr %MyStruct*, %MyStruct* @a_gv, i32 0, i32 2
```

GEPs with structs




```
%MyStruct = type < { i8, i32, [3 x i32] }>
```

```
@a_gv = global %MyStruct { i8 99, i32 17, [3 x i32] [i32 1, i32 2, i32 3] }
```

```
%new_ptr = getelementptr %MyStruct*, %MyStruct* @a_gv, i32 0, i32 2, i32 0
```

GEPs with structs



```
%MyStruct = type < { i8, i32, [3 x i32] }>
```

```
@a_gv = global %MyStruct { i8 99, i32 17, [3 x i32] [i32 1, i32 2, i32 3] }
```

```
%new_ptr = getelementptr %MyStruct*, %MyStruct* @a_gv, i32 0, i32 2, i32 1
```


GEP fundamentals

1. Understand the first index:

- It does NOT change the pointer type.
- It offsets by the **pointee type**.


2. Further indices:

- Offset inside **aggregate types**.
- Change the pointer type by removing one layer of “aggregation”.

3. Struct indices must be constants.

GEPs with structs

Goal: get an i8* to 2nd element of @gv2



```
@gv1 = global [4 x i8] { [i8 0, i8 1, i8 2, i8 3] }
```

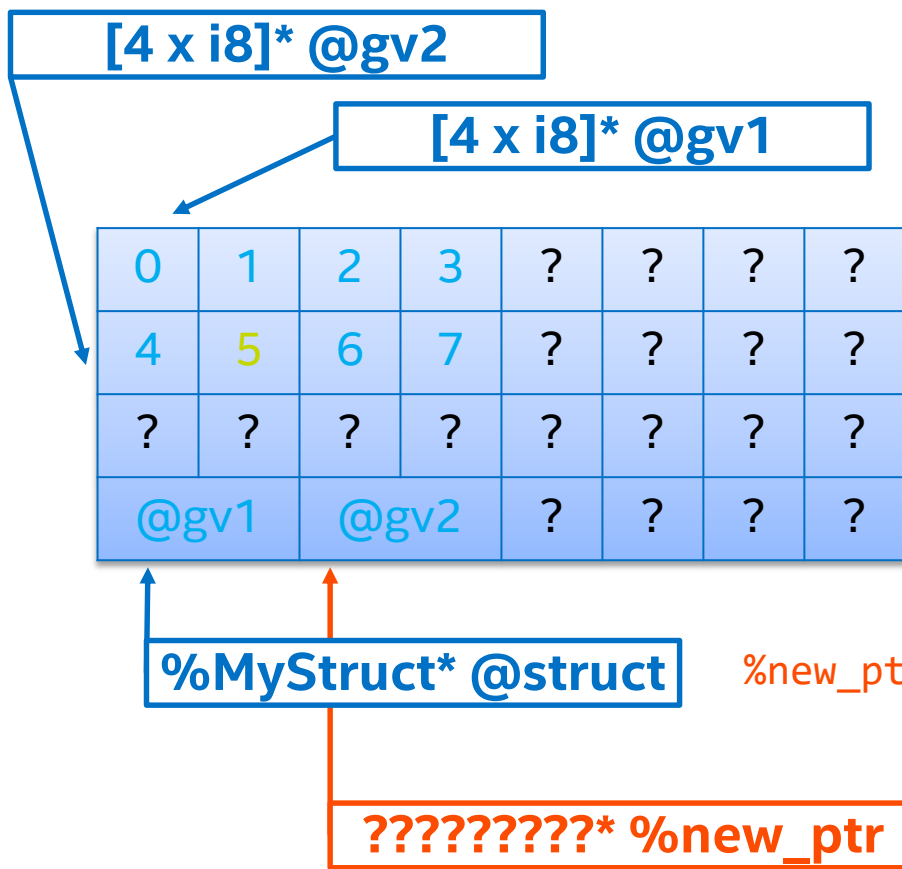
```
@gv2 = global [4 x i8] { [i8 4, i8 5, i8 6, i8 7] }
```

```
%MyStruct = type { [4 x i8]*, [4 x i8]* }
```

```
@struct = global %MyStruct { [4 x i8]* @gv1, [4 x i8]* @gv2 }
```

GEPs with structs

Goal: get an i8* to 2nd element of @gv2 using @struct



```
@gv1 = global [4 x i8] { [i8 0, i8 1, i8 2, i8 3] }
```

```
@gv2 = global [4 x i8] { [i8 4, i8 5, i8 6, i8 7] }
```

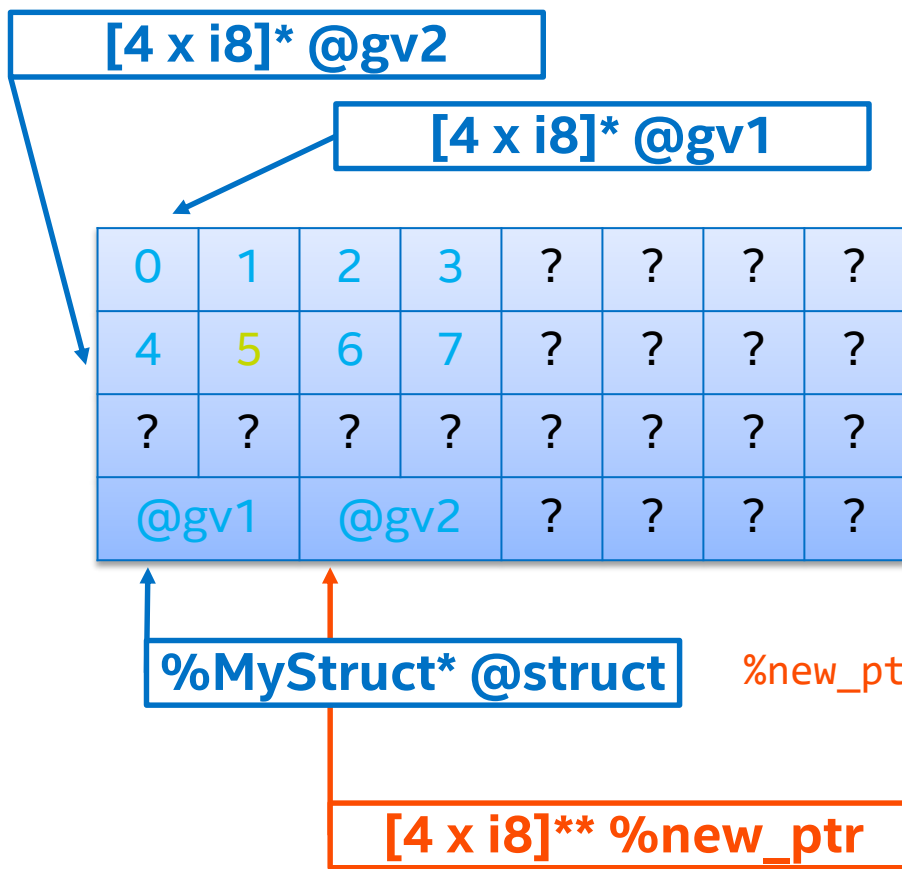
```
%MyStruct = type { [4 x i8]*, [4 x i8]* }
```

```
@struct = global %MyStruct { [4 x i8]* @gv1, [4 x i8]* @gv2 }
```

```
%new_ptr = getelementptr %MyStruct, %MyStruct* @struct, i32 0, i32 1
```

GEPs with structs

Goal: get an `i8*` to 2nd element of `@gv2` using `@struct`



```
@gv1 = global [4 x i8] { [i8 0, i8 1, i8 2, i8 3] }
```

```
@gv2 = global [4 x i8] { [i8 4, i8 5, i8 6, i8 7] }
```

```
%MyStruct = type { [4 x i8]*, [4 x i8]* }
```

```
@struct = global %MyStruct { [4 x i8]* @gv1, [4 x i8]* @gv2 }
```

```
%new_ptr = getelementptr %MyStruct, %MyStruct* @struct, i32 0, i32 1
```

Quiz: Can we add an extra zero to this GEP?

GEP fundamentals

1. Understand the first index:

- It does NOT change the pointer type.
- It offsets by the **pointee type**.

2. Further indices:

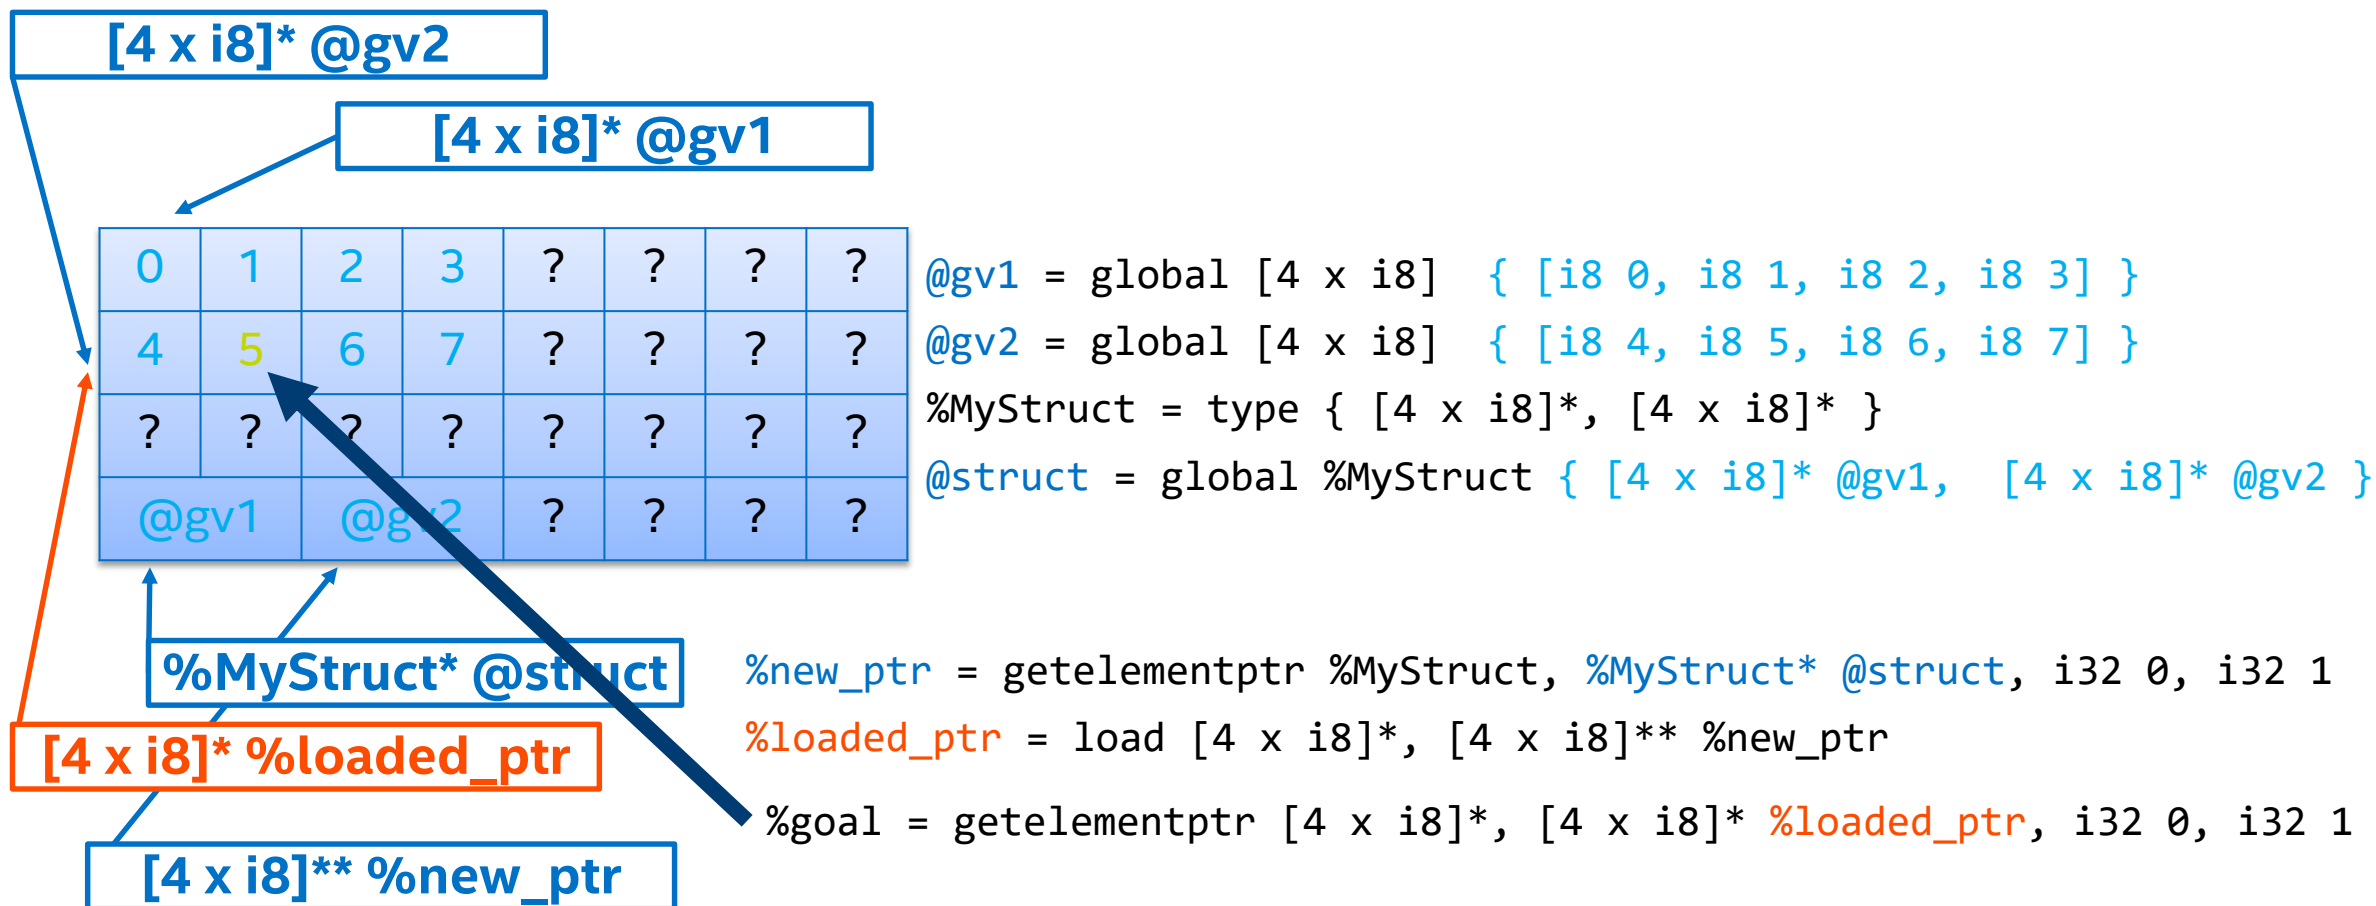
- Offset inside **aggregate types**.
- Change the pointer type by removing one layer of “aggregation”.

3. Struct indices must be constants.

4. GEPs never load from memory!

GEPs with structs

Goal: get an i8* to 2nd element of @gv2 using @struct



Final remarks:

LLVM IR is constantly evolving.

Covered fundamental topics unlikely to change soon.

- There is a lot more to explore!

[Some] Topics not covered:

1. Constants
2. Constant expressions
3. Intrinsics
4. Exceptions
5. Debug information
6. Metadata
7. Attributes
8. Vector instructions

Final remarks:

LLVM IR is constantly evolving.

Covered fundamental topics unlikely to change soon.

- There is a lot more to explore!

Next steps:

- Learn how to manipulate IR using the LLVM library
- Look at the OPT code!

THANK YOU!

QUESTIONS?

Disclaimer

Intel, the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries.

***Other names and brands may be claimed as the property of others.**

© Intel Corporation