

Outline

- Navigating and changing the IR
- The machine type system
- Using and writing passes

LLVM, in Greater Detail

Thanks to Gabe Weisz for some content.

1

2

LLVM Overview

- C++-based compiler framework
- (Fairly) well documented API
- Structures to help you process programs
 - Iterators for modules, functions, blocks, uses
 - Functions to inspect data types and constants
 - Many classes have dump() member functions that print instances to standard error
 - In GDB, use `p obj->dump()` to see the contents of obj

3

LLVM IR

- Mostly machine-independent assembly
 - Target triples define alignment, pointer sizes
- Arbitrary number of “registers”
 - Really, stack locations or SSA values
 - Virtual registers appear in lower-level IRs
- Locals start with %, globals with @
 - Instructions that produce values can be named

4

Iterators – navigating the IR

- `Module::iterator`
 - Modules are translation units
 - Iterates through functions in the module
- `Function::iterator`
 - Iterates through a function's basic blocks
- `BasicBlock::iterator`
 - Iterates through instructions in a block
- `Value::use_iterator`
 - Iterates through *uses*
 - Instructions are values; so are constants
 - How does SSA help?
- `User::op_iterator`
 - Iterates over operands. (Instructions are users!)
 - Many instruction classes define convenient accessors
 - `LoadInst::getPointerOperand`

5

Instructions

- Instruction is subclassed for various types of operation
 - `LoadInst`, `StoreInst`, `CmpInst`, `BranchInst`, etc
- Most arithmetic operations are `BinaryOperators` that contain a code for the operation
- Some instructions can only appear in certain places
 - Branches are only at the end of a basic block
 - Phi instructions are only at the beginning

7

More Iterators

- Some iterators wrap other iterators
 - `inst_iterator` walks over all instructions in a function
 - `for (inst_iterator II = inst_begin(f), IE = inst_end(f); II != IE; ++II)`
 - In `Transforms/Utils/FunctionUtils.h`
- Prefer `++i` over `i++` and precompute the end iterator
 - Especially for fancier iterators
- Most iterators automatically cast to a pointer to the object type (`inst_iterator` does not)
- Be careful about modifying the object you're iterating over during iteration

6

(Re)moving Instructions

- `eraseFromParent()`
 - Remove from basic block, drop all references, delete
- `removeFromParent()`
 - Remove from basic block
 - Use if you will re-attach the instruction
 - Does not drop references (or clear the use list), so if you don't re-attach it Bad Things will happen
- `moveBefore/insertBefore/insertAfter` are also available
- `ReplaceInstWithValue` and `ReplaceInstWithInst` are also useful to have

8

IR Types

- Primitive types
 - Integers (iN of N bits, N from 1 to $2^{23}-1$)
 - Floating point (half, float, double, x86_fp80, ...)
 - Weirdos (x86mmx, void, ...)
- Derived types
 - Arrays ($[\# \text{elements} (\geq 0) \times \text{element type}]$)
 - Functions ($\text{return type} (\text{paramlist})$)
 - Pointers ($\text{type}^*, \text{type addr space}(N)^*$)
 - Vectors ($\langle \# \text{elements} (> 0) \times \text{element type} \rangle$)
 - Structures ($\{ \text{typelist} \}$) ...

9

IR Types

- getelementptr instruction gives you the address of a field or an array cell (why have this?)
- Types in the LLVM IR are structural
 - Mostly compared by shape, not by name
 - Once allocated, there is only one 32-bit integer until the end of time
 - Only one instance of a given shape exists at once
 - Benefit?
 - Exception: “identified” structures
 - Problem: how do you write down the type of a singly-linked list?

10

Identified Structs

- LLVM 3 came with a redesign of the IR type system mainly over the issue of recursive and abstract types
- Literal structs are compared by shape and must not be recursive; all fields must be known
 - $\{i32, i32\}$; StructType::get
- Only identified structs can be recursive; declaration and definition of fields is separate
 - $\%T1 = \text{type} \{ \text{type list} \}$; StructType::create
 - type list may directly or indirectly refer to $\%T1$
 - $\%T1 = \text{type opaque}$ - “I’ll fill in the fields later.”
- LLVM may rename your identifiers

11

Examples of Types

```
[40 x i32]
[3 x [4 x i32]]
%sLL =
    type { i32, %sLL* }
i32 (%sLL*)*

%struct.a =
    type { %struct.b* }
%struct.b =
    type opaque

{ i32, [0 x float] }
```

12

Last words about opaque types

- Opaque types are not void* (C void* is i8*)
- Consider `%local_a = alloca %a, align 4`
 - `%a = type { %b* }`
`%b = type opaque`
 - OK!
 - `%a = type { %b }`
`%b = type opaque`
 - “Cannot allocate unsized type”
 - `%a = type { %b }`
`%b = type { %a }`
 - “Segmentation fault”

13

Module Verifier (-verify)

- `opt` runs this automatically unless you disable it
- Sanity-checks passes
 - You may need to break module invariants while operating on them, eg:
 - Types of binary operator parameters are the same
 - Terminators (branches) only at the end of basic blocks
 - Functions are called with correct argument types
 - Instructions belong to basic blocks
 - Constants in a switch are the right type
 - Entry node has no predecessors (*and so on...*)

15

Passes

- For assignments, don't use provided LLVM passes unless instructed to
 - We want you to implement them yourself to understand how they really work
- For projects, use whatever you want
- Two major kinds of passes
 - Analysis: provide information
 - Transform: modify the program

14

The mem2reg transform pass

- The LLVM IR is natively SSA
 - An Instruction is the same thing as the Value it produces
 - `%foo = inst` in the LLVM IR just gives a name to inst in the syntax; **%foo does not exist inside the compiler**
- It may be nontrivial for frontends to emit SSA directly
- `mem2reg` understands certain use patterns that frontends use to emit “variables”

16

mem2reg conventions

```
int ssa1() {
    int z = f() + 1;
    return z;
}
```

```
define i32 @ssa1() nounwind {
entry:
    %call = call i32 @f()
    %add = add nsw i32 %call, 1
    ret i32 %add
}
```

alloca in the entry block

only used by load and store

```
define i32 @ssa1() nounwind {
entry:
    %z = alloca i32, align 4
    %call = call i32 @f()
    %add = add nsw i32 %call, 1
    store i32 %add, i32* %z, align 4
    %0 = load i32* %z, align 4
    ret i32 %0
}
```

17

mem2reg might add SSA features

```
int ssa2() {
    int y, z;
    y = f();
    if (y < 0)
        z = y + 1;
    else
        z = y + 2;
    return z;
}
```

```
define i32 @ssa2() nounwind {
entry:
    %call = call i32 @f()
    %cmp = icmp slt i32 %call, 0
    br i1 %cmp, label %if.then, label %if.else
```

```
if.then:
    %add = add nsw i32 %call, 1
    br label %if.end
```

```
if.else:
    %add1 = add nsw i32 %call, 2
    br label %if.end
```

```
if.end:
    %z.0 = phi i32 [ %add, %if.then ], [ %add1, %if.else ]
    ret i32 %z.0
}
```

18

Rules for Phi instructions

- phi *type* [*val1*, *inedge1*], [*val2*, *inedge2*]
 - Select *val1* if coming from *inedge1*; *val2* if from *inedge2*
- Phi nodes may refer to themselves (loops!) and may select undef (undefined) values for certain in-edges
- Placement requirements:
 - must be at the beginning of the block
 - must have one entry for each predecessor
 - must have at least one entry

mem2reg confuses easily

```
int ssa3() {
    int z;
    return *(&z + 1 - 1);
}
```

```
define i32 @ssa3() nounwind {
entry:
    %z = alloca i32, align 4
    %add.ptr = getelementptr inbounds i32* %z, i32 1
    %add.ptr1 = getelementptr inbounds i32* %add.ptr, i32 -1
    %0 = load i32* %add.ptr1, align 4
    ret i32 %0
}
```

getelementptr
abstracts away
offset calculation

Why not make mem2reg smarter?

(note that compiling with -O2 optimizes this down to ret undef)

19

20

Loop information (-loops)

- Analysis/LoopInfo.h
- Basic blocks in a loop
- Headers and pre-headers
- Exit and exiting blocks
- Back-edges
- “Canonical induction variable”
 - Starts at 0 and counts up by 1?
 - Starts at some number and counts down to 0
- Loop count

21

Using an analysis pass

- `#include "llvm/Analysis/LoopInfo.h"`
- `AU.addRequired<LoopInfo>();`
in `getAnalysisUsage()`
- `LoopInfo& LI = getAnalysis<LoopInfo>(F);`
in a function called from `runOnModule`
with function `F` inside that module
- `LI.dump();`
“Loop at depth 1 containing:
%for.cond<header><exiting>,%for.body,
%for.inc<latch>” (from `loop.c`)
- `PassManager` sequences both kinds of passes using `getAnalysisUsage()`

22

Scalar Evolution (-scalar-evolution)

- Tracks changes to variables through multiple loop nests
- Gives start value, step size, kind of evolution
 - Constant
 - Add a value each iteration
 - Multiply a value each iteration
 - More complicated relationships as well
- Useful to aggregate accesses into arrays into larger blocks or to improve cache performance

23

Target Data (-targetdata)

- Endian-ness
- Pointer sizes
- Alignment
- Actual size (in bits) of variables
- Actual layout of structures
 - (taking into account platform alignment requirements)

24

Alias Analyses

```
%1 = load i32* %A
store i32 5, i32* %B
%3 = add i32 %1, i32 9
store i32 %3, i32* %C
```


If we know that
%A != %B != %C
we have more freedom to
reorder code,
promote to registers, etc.

- LLVM includes a number of passes that collect various kinds of alias information
- Can get information about both global and local variables
- Included passes take into account the behavior of the C standard library (eg, `sin()` will not make new aliases)

25

Simplify CFG

- A cleanup pass
- Removes unnecessary basic blocks by merging unconditional branches if the second block has only one predecessor



- Removes unreachable blocks
- Removes Phi instructions in blocks with single predecessors

26

Other useful passes

- Liveness-based dead code elimination
 - Assumes code is dead unless proven otherwise
- Sparse conditional constant propagation
 - Aggressively search for constants
- Correlated propagation
 - Replace select instructions that select on a constant
- Loop invariant code motion
 - Move code out of loops where possible
- Dead global elimination
- Canonicalize induction variables
 - All loops start from 0
- Canonicalize loops
 - Put loop structures in standard form

27

Notes on Writing Passes


- Declare which passes you use (and what your pass mutates) in `getAnalysisUsage`
- The `CommandLine` library allows you to add command line parameters very quickly
 - Conflicts in parameter names won't show up until runtime, since passes are loaded dynamically
- Be mindful of correctness; the module verifier is like a syntax checker
 - Does your pass make sense in a multithreaded environment?

28

Transformations and memory

- For regular loads/stores, LLVM forbids introducing new stores to externally observable locations:

```
int x;
void f(int* a) {
    for (int i = 0; i < 100; i++) {
        if (a[i])
            x += 1;
    }
}
```



```
int x;
void f(int* a) {
    int xtemp = x;
    for (int i = 0; i < 100; i++) {
        if (a[i])
            xtemp += 1;
    }
    x = xtemp;
}
```

- `volatile` marks memory operations that cannot be reordered (wrt `volatile` operations) or removed
- To support new features in C++11, LLVM provides other *atomic orderings* that can be applied to loads and stores

29

Links

- When in doubt, read the documentation—and the code!
 - <http://llvm.org/doxygen/>
- Articles on the LLVM site are very useful
 - <http://llvm.org/docs/Passes.html>
 - <http://llvm.org/docs/ProgrammersManual.html>
 - <http://llvm.org/docs/Atomics.html>
 - <http://llvm.org/docs/LangRef.html>

31

Projects using LLVM

- Just a few from llvm.org:
 - Clang: a C-family language frontend
 - LLDB: an improved debugger using Clang data
 - vmkit: building Java/.NET VMs
 - klee: a symbolic virtual machine for LLVM IR
- Emscripten: LLVM bitcode → JavaScript
 - <http://emscripten.org/>
- Vellvm: a formalization of the LLVM IR
 - <http://www.cis.upenn.edu/~jianzhou/Vellvm/>

30

Supplemental

32

High-level view of atomic orderings

- NotAtomic – ordinary loads/stores; races are undefined
- Unordered – races have “somewhat sane” results
 - A load cannot see a value which was never stored
 - May not be supported for all types on all platforms
 - Used for shared variables in Java, “safe” languages
- Monotonic – single locations have consistent order
- Acquire/Release – when paired together, strong enough to write a lock
- SequentiallyConsistent – Acquire loads, Release stores, and all SequentiallyConsistent operations have a total order
 - Java volatile

33


Atomic orderings and you

- If your project is about fine-grained parallelism or lock-free data structures, you need to think about these things
 - LLVM also exposes cmpxchg, atomicrmw, fence
 - LLVM does not expose LL/SC
- If not, just don't, in general:
 - reorder LLVM volatile operations wrt other volatile operations
 - introduce new stores to (shared) locations that would not have been previously stored to
- More details at <http://llvm.org/docs/LangRef.html#memmodel>

34

Your module, post-IL

- LLVM still has to generate machine code!
- Your module goes through ~3 more stages:



35

LLVM IL → SelectionDAG

```
define i32 @foo(i32 %a, i32 %b) nounwind {  
entry:  
  %cmp = icmp sgt i32 %a, %b  
  br i1 %cmp, label %if.then, label %if.else
```

```
if.then:  
  %add = add nsw i32 %a, %b  
  br label %return
```

```
if.else:  
  %sub = sub nsw i32 %a, %b  
  br label %return
```

```
return:  
  %retval.0 = phi i32 [ %add, %if.then ],  
    [ %sub, %if.else ]  
  ret i32 %retval.0
```


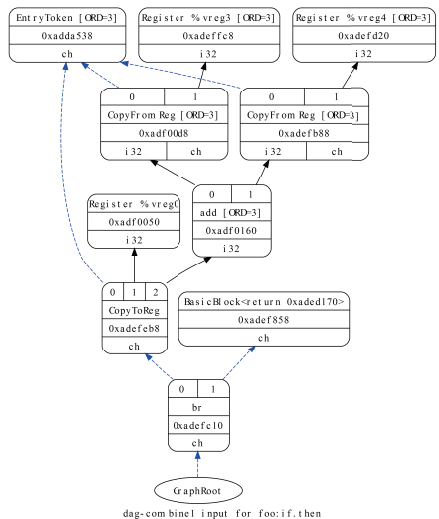
```
}
```

36

LLVM IL \rightarrow SelectionDAG

```
if.then:
%add = add nsw i32 %a, %b
br label %return
```

```
return:
%retval.0 = phi i32 ...
ret i32 %retval.0
```



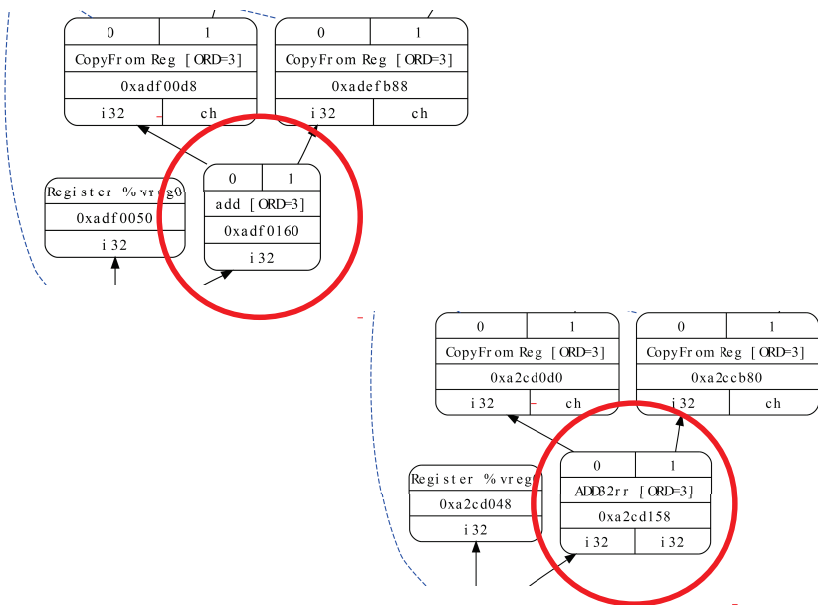
37

SelectionDAG transformations

- Build initial DAG from LLVM IR
- Simplify!
- Legalize types (vectors \rightarrow scalars)
- Simplify!
- Legalize ops (x86 doesn't do byte-size CMOVs)
- Simplify!
- Select instructions
- Schedule instructions

38


Instruction selection



39

Target-Independent Code Generation

- There are a lot of different machines!
 - And even on x86, there are a lot of different ADDs!
- Make the process data-driven



40

tablegen

- The tablegen tool is run when compiling the LLVM library for each target
- It accepts a custom text-based record description format and generates C++ definitions using various backends
 - defm ADD : ArithBinOp_RF<0x00, 0x02, 0x04, "add", MRM0r, MRM0m, X86add_flag, add, 1, 1>;
 - ArithBinOp_RF is actually another macro...
- There is still a lot of human-written code in the backends (X86 instruction encoding, for one)

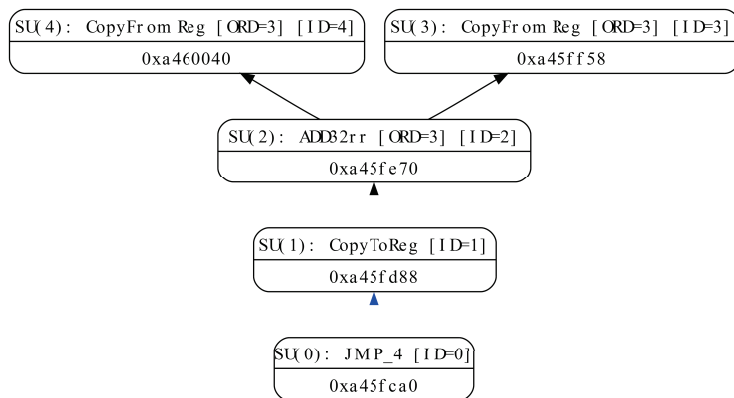
ADD32rr

```
def ADD32rr {
  Format BinOpRR:f = MRMDestReg;
  Domain I:d = GenericDomain;
  string Namespace = "X86";
  dag OutOperandList = (outs GR32:$dst);
  dag InOperandList = (ins GR32:$src1, GR32:$src2);
  string AsmString = "add{l} {$src2, $src1|$src1, $src2}";
  list<dag> Pattern =
    [(set GR32:$dst, EFLAGS,
      (X86add_flag GR32:$src1, GR32:$src2))];
  list<Register> Uses = [];
  list<Register> Defs = [EFLAGS];
  list<Predicate> Predicates = [];
  InstrItinClass Itinerary = IIC_BIN_NONMEM;
  string Constraints = "$src1 = $dst";
  bits<8> Opcode = { 0, 0, 0, 0, 0, 0, 0, 1 };
  Format Form = MRMDestReg;
  bits<6> FormBits = { 0, 0, 0, 0, 1, 1 };
  ... lots more ...
}
```

41

42

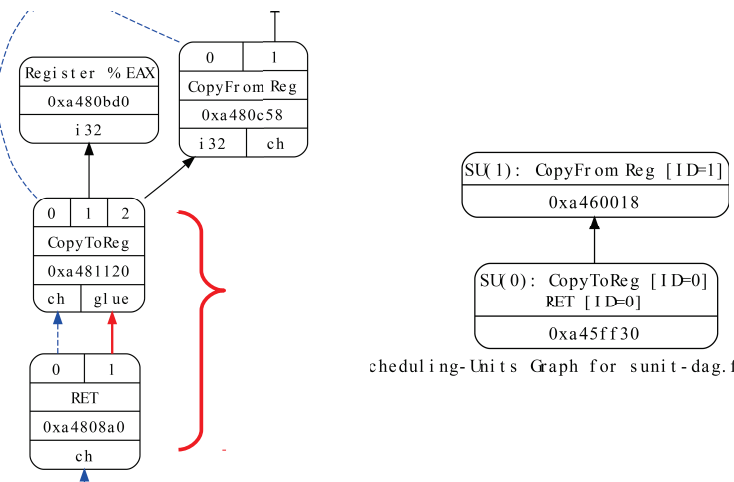
Scheduling



Scheduling-Units Graph for sunit-dag.foo:if.then

We have choices to make
(here, which CopyFromReg executes first?)

Scheduling



scheduling-Units Graph for sunit-dag.foo:retur

Chains add control dependency.
Glue forbids breaking up instructions.

43

44

Lowering to MC

BB#1: derived from LLVM BB %if.then

```
%vreg0<def,tied1> = ADD32rr  
  %vreg3<tied0>,  
  %vreg4,  
  %EFLAGS<imp-def,dead>;  
  GR32:%vreg0,%vreg3,%vreg4  
JMP_4 <BB#3>
```

Virtual registers

Tied registers

Register classes

BB#3: derived from LLVM BB %return

```
%vreg2<def> = PHI  
  %vreg1, <BB#2>, %vreg0, <BB#1>;  
  GR32:%vreg2,%vreg1,%vreg0  
%EAX<def> = COPY %vreg2; GR32:%vreg2  
RET
```

Still in SSA

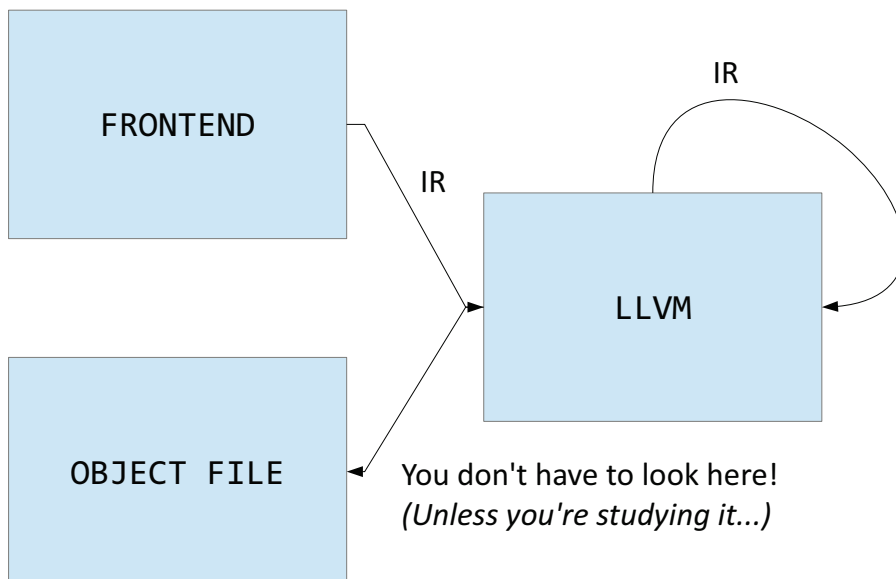
45

MC transformations

- SSA-based MC optimizations
- Register allocation
 - 2AC correction and **Leave SSA**; copy coalescing; add spillcode
- Prolog/epilog insertion
- Stack layout
- Last-chance MC optimizations/spillcode scheduling
- Encoding

46

Thanks, abstraction!



47