



Z-Buffer Optimizations

Patrick Cozzi

Analytical Graphics, Inc.

Overview



Z-Buffer Review



Hardware: Early-Z



Software: Front-to-Back Sorting



Hardware: Double-Speed Z-Only



Software: Early-Z Pass



Software: Deferred Shading



Hardware: Buffer Compression



Hardware: Fast Clear



Hardware: Z-Cull



Future: Programmable Culling Unit

Z-Buffer Review



- Also called Depth Buffer
- Fragment vs Pixel
- Alternatives: Painter's, Ray Casting, etc




Z-Buffer History

- “Brute-force approach”
- “Ridiculously expensive”
- Sutherland, Sproull, and, Schumacker, “A Characterization of Ten Hidden-Surface Algorithms”, 1974



Z-Buffer Quiz

- 10 triangles cover a pixel. Rendering these in random order with a Z-buffer, what is the average number of times the pixel's z-value is written?





Z-Buffer Quiz

- 1st triangle writes depth
- 2nd triangle has 1/2 chance of writing depth
- 3rd triangle has 1/3 chance of writing depth
- $1 + 1/2 + 1/3 + \dots + 1/10 = 2.9289\dots$



Z-Buffer Quiz


Harmonic Series

# Triangles	# Depth Writes
1	1
4	2.08
11	3.02
31	4.03
83	5
12,367	10




Z-Test in the Pipeline



- When is the Z-Test?



or





Early-Z



- Avoid expensive fragment shaders
- Reduce bandwidth to frame buffer
 - Writes not reads

Early-Z




- Automatically enabled on GeForce (8?) unless
 - Fragment shader **discards** or write depth
 - Depth writes and alpha-test are enabled
- Fine-grained as opposed to *Z-Cull*.
- ATI: “Top of the Pipe Z Reject”

See NVIDIA GPU Programming Guide for exact details

Front-to-Back Sorting




- Utilize *Early-Z* for opaque objects
- Old hardware still has less z-buffer writes
- CPU overhead. Need efficient sorting
 - Bucket Sort
 - Octtree
- Conflicts with state sorting



1	2
0	1

Double Speed Z-Only




- GeForce FX and later render at double speed when writing only depth or stencil
- Enabled when
 - Color writes are disabled
 - Fragment shader **discards** or write depth
 - Alpha-test is disabled

See NVIDIA GPU Programming Guide for exact details

Early-Z Pass



- Software technique to utilize *Early-Z* and *Double Speed Z-Only*
- Two passes
 - Render depth only. “Lay down depth” – *Double Speed Z-Only*
 - Render with full shaders – *Early-Z* (and *Z-Cull*)





Deferred Shading



- Similar to *Early-Z Pass*
 - 1st Pass: Visibility tests
 - 2nd Pass: Shading
- Different than *Early-Z Pass*
 - Geometry is only transformed once

Deferred Shading




○ 1st Pass


- Render geometry into *G-Buffers*:



Fragment Colors



Normals



Depth



Edge Weight

Images from Tabula Rasa. See Resources.



Deferred Shading



○ 2nd Pass

- Shading == post processing effects
- Render full screen quads that read from *G-Buffers*
- Objects are no longer needed

Deferred Shading



- Light Accumulation Result



Image from Tabula Rasa. See Resources.




Deferred Shading




- Eliminates shading fragments that fail Z-Test
- Increases video memory requirement
- How does it affect bandwidth?

Buffer Compression



- Reduce depth buffer bandwidth
- Generally does not reduce memory usage of actual depth buffer
- Same architecture applies to other buffers, e.g. color and stencil

Buffer Compression



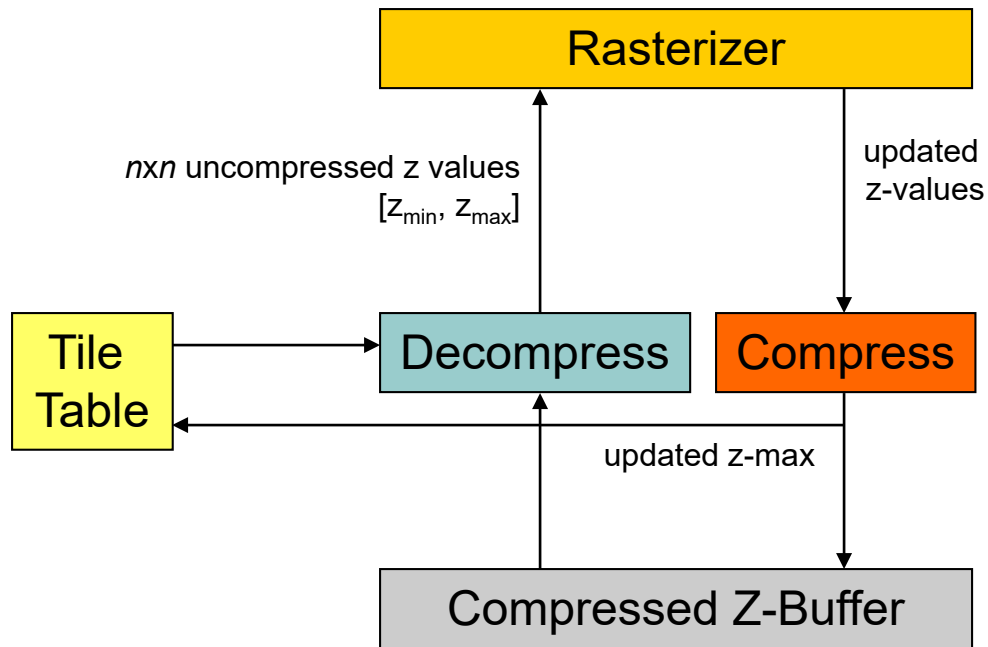

- Tile Table: Status for $n \times n$ tile of depths, e.g. $n=8$
 - [state, z_{\min} , z_{\max}]
 - state is either *compressed*, *uncompressed*, or *cleared*

[*uncompressed*, 0.1, 0.8]




0.1	0.5	0.5	0.1
0.5	0.8	0.8	0.5
0.5	0.8	0.8	0.5
0.1	0.5	0.5	0.1

Buffer Compression




Buffer Compression



○ Depth Buffer Write

- Rasterizer modifies copy of uncompressed tile
- Tile is lossless compressed (if possible) and sent to actual depth buffer
- Update Tile Table
 - z_{\min} and z_{\max}
 - status: *compressed* or *decompressed*

Buffer Compression




- Depth Buffer Read

- Tile Status


- *Uncompressed*: Send tile
 - *Decompress*: Decompress and send tile
 - *Cleared*: See Fast Clear

Fast Clear




- Don't touch depth buffer
- `glClear` sets state of each tile to *cleared*
- When the rasterizer reads a cleared buffer
 - A tile filled with `GL_DEPTH_CLEAR_VALUE` is sent
 - Depth buffer is not accessed

Fast Clear

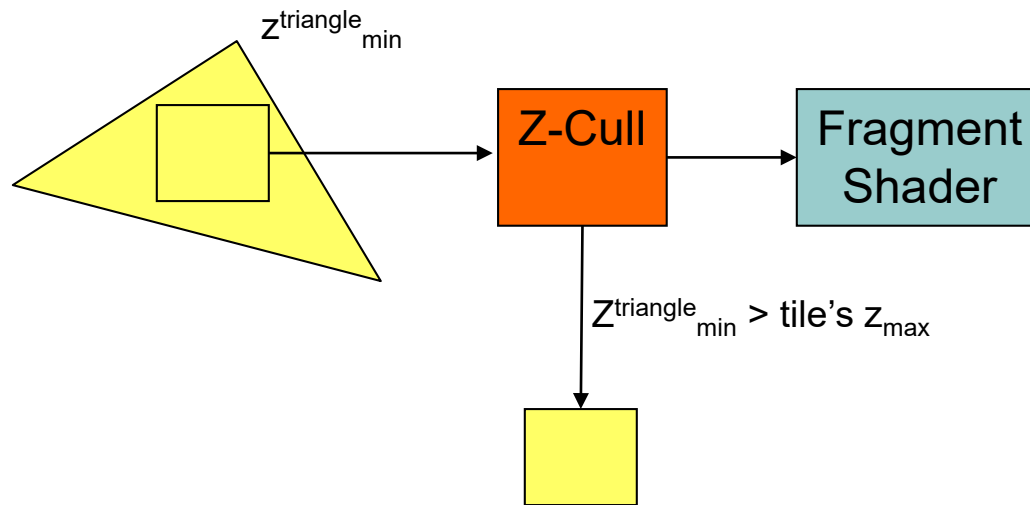


- Use `glClear`
 - Not full screen quads
 - No "one frame positive, one frame negative" trick
- Clear stencil together with depth


Z-Cull



- Cull blocks of fragments before shading
- Coarse-grained as opposed to *Early-Z*

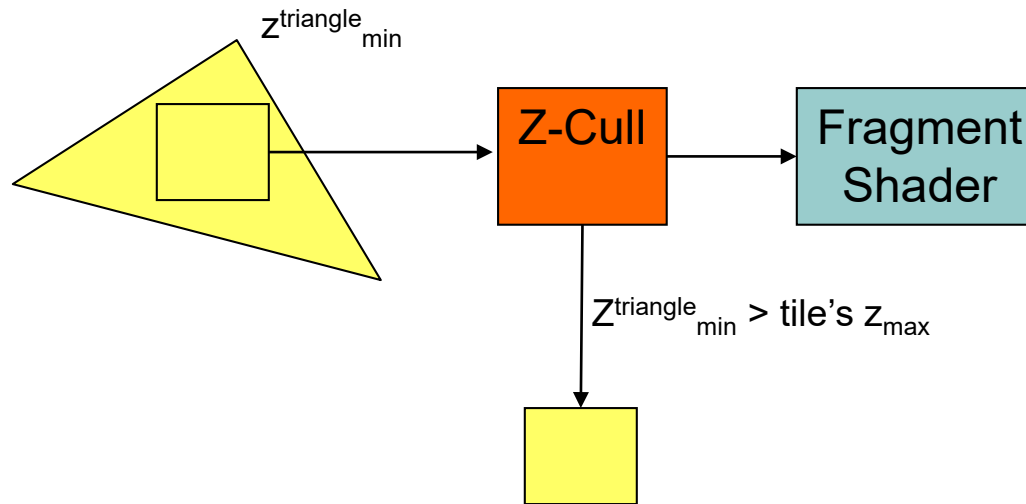


Z-Cull




○ Z_{\max} -Culling

- Rasterizer fetches z_{\max} for each tile it processes
- Compute $z_{\min}^{\text{triangle}}$ for a triangle
- Culled if $z_{\min}^{\text{triangle}} > z_{\max}$

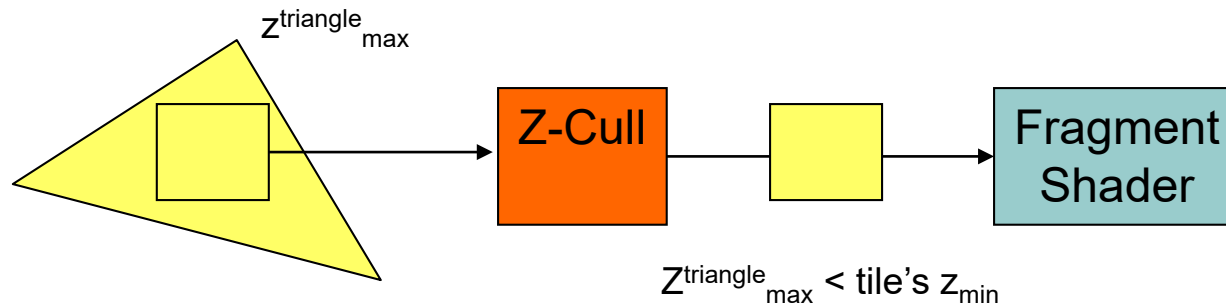


Z-Cull




○ Z_{\min} -Culling

- Support different depth tests
- Avoid depth buffer reads
- If triangle is in front of tile, depth tests for each pixel is unnecessary



Z-Cull



- Automatically enabled on GeForce (6?) cards unless
 - `glClear` isn't used
 - Fragment shader writes depth (or `discards`?)
 - Direction of depth test is changed
- ATI recommends avoiding `=` and `!=` depth compares and stencil fail and stencil depth fail operations
- Less efficient when depth varies a lot within a few pixels

See NVIDIA GPU Programming Guide for exact details



Programmable Culling Unit




- Cull before fragment shader even if the shader writes depth or **discards**
- Run part of shader over an entire tile to determine lower bound z value
- Hasselgren and Akenine-Möller, “PCU: The Programmable Culling Unit,” 2007

Summary

- What was once “ridiculously expensive” is now the primary visible surface algorithm for rasterization



Resources



Sections 7.9.2 and 18.3

www.realtimerendering.com

Resources



GeForce 8 Guide: sections 3.4.9, 3.6, and 4.8

GeForce 7 Guide: section 3.6

developer.nvidia.com/object/gpu_programming_guide.html

Resources



ATI Radeon HyperZ Technology
Steve Morein

Resources




Performance Optimization Techniques for ATI
Graphics Hardware with DirectX® 9.0

Guennadi Riguer

Sections 6.5 and 8

http://ati.amd.com/developer/dx9/ATI-DX9_Optimization.pdf

Resources



Chapter 28: Graphics Pipeline Performance

developer.nvidia.com/object/gpu_gems_home.html

Resources



Chapter 19: Deferred Shading in Tabula Rasa

developer.nvidia.com/object/gpu-gems-3.html